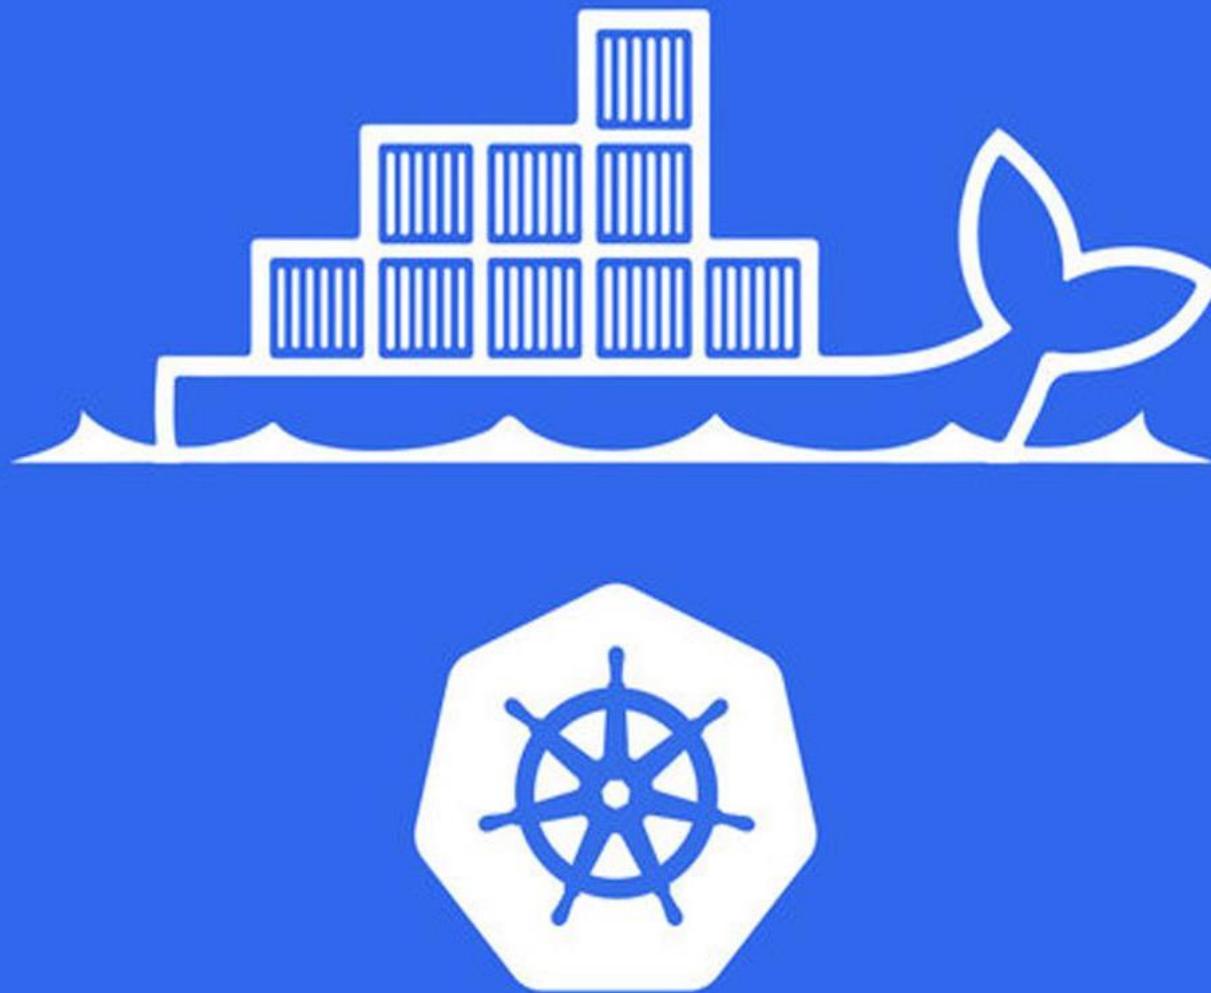


# KUBERNETES ADMINISTRATOR



Gaurav

RHCSA | RHCE | CKA | SCA | SCE | SCA+ HA | ANSIBLE | DOCKER | OPENSHIFT® | OPEN SOURCE TECHNOLOGIES

# TABLE OF CONTENT

- Core Concepts
- Installation of Kubernetes Cluster
- Using Kubernetes Features
- Networking in Kubernetes
- Security in Kubernetes
- Storage
- Logging & Monitoring

# WHAT IS CONTAINER?

Containers are a way for developers to easily package and deliver applications, and for operations to easily run them anywhere in seconds, with no installation or setup necessary

# WHY CONTAINER?

Containers require less system resources than traditional or hardware virtual machine environments because they don't include operating system images. Increased portability. Applications running in containers can be deployed easily to multiple different operating systems and hardware platforms.

# CONTAINER ORCHESTRATION

# SCENARIO

# SCENARIO OF GUARD

- You are owner of a building and you have 5 spots where people can enter your building and you want 5 security guards guarding the spots. All good till now.
- Now consider one of the guard was out of service for 2 hours due to some personal reasons. Now as a building owner its your responsibility to guard or employ another guard replacing the existing. Do you like to be manually interrupted from your task to look after who is out and whom to replace.
- No, no one likes to be. Now the solution could be, go to a third party vendor who provides 24\*7 availability of the guards. Its the responsibility of the vendor to make 24\*7 availability based on the requirements.

# WHAT PROBLEM DOES KUBERNETES SOLVE?

Now lets replace the characters.

- Building -> Your application
- Owner -> The application owner
- Guards -> Containers
- Vendor -> Kubernetes
- Configuration (repliica)-> That we make a deal with vendor to maintain appropriate conf 24\*7.

# CONTAINER ORCHESTRATION

- Container orchestration automates the deployment, management, scaling, and networking of containers. Enterprises that need to deploy and manage hundreds or thousands of Linux® containers so in this case we don't want downtime so container orchestration will give this facility.

# What is container orchestration used for?

- Provisioning and deployment
- Configuration and scheduling
- Resource allocation
- Container availability
- Scaling or removing containers based on balancing workloads across your infrastructure
- Load balancing and traffic routing
- Monitoring container health
- Configuring applications based on the container in which they will run
- Keeping interactions between containers secure

# TOOLS

- DOCKER SWARM
- KUBERNETES

# DIFFERENCE BETWEEN DOCKER SWARM AND KUBERNETES

# INSTALLATION CLUSTER CONFIGURATION

## DOCKER SWARM

- SETTING UP THE CLUSTER IS SIMPLE
- REQUIRES ONLY 2 COMMANDS

## KUBERNETES

- SETTING UP THE CLUSTER IS CHALLENGING AND COMPLICATED

# USER INTERFACE

## DOCKER SWARM

- THERE IS NOT GUI AVAILABLE

## KUBERNETES

- PROVIDES A GUI  
(KUBERNETES DASHBOARD)

# SCALABILITY

## DOCKER SWARM

- Scaling up is 5X faster than kubernetes
- Better when 10-20 containers

## KUBERNETES

- Scaling up is easy
- Better when 100-1000 containers in pod

# AUTO SCALING

## DOCKER SWARM

- Scaling up or scaling down has to be done manually

## KUBERNETES

- Based on server traffic, containers will be scaled automatically

# LOAD BALANCING

## DOCKER SWARM

- SWARM DOES AUTO LOAD BALANCING

## KUBERNETES

- Manual configuration needed for load balancing traffic

# ROLLING UPDATE AND ROLLBACKS

## DOCKER SWARM

- Rolling updates progressively updates the containers one after the other while ensuring HA
- No automatic Rollbacks

## KUBERNETES

- Rolling updates progressively updates the pods one after the other while ensuring HA
- Automatic Rollbacks in case of failure

# DATA VOLUMES

## DOCKER SWARM

- Storage volumes can be shared with any other container in the node

## KUBERNETES

- Storage volumes can be shared only between containers within the same pod

# LOGGING AND MONITORING

## DOCKER SWARM

- 3rd party logging and monitoring tools required

## KUBERNETES

- In-built logging and monitoring tools in place

# WHAT IS KUBERNETES

Kubernetes is a system for managing containerized applications across a cluster of nodes.

# Why you need Kubernetes and what can it do?

Containers are a good way to bundle and run your applications. In a production environment, you need to manage the containers that run the applications and ensure that there is no downtime. For example, if a container goes down, another container needs to start. Wouldn't it be easier if this behavior was handled by a system?

**That's how Kubernetes comes to the rescue!**

# KUBERNETES PROVIDES

- Service Discovery and Load Balancing
- Storage Orchestration
- Automated rollouts and rollback
- Automatic provides nodes as per container requirement
- Self-healing
- Secret and configuration management



# kubernetes



Self-Healing



Automated Rollbacks



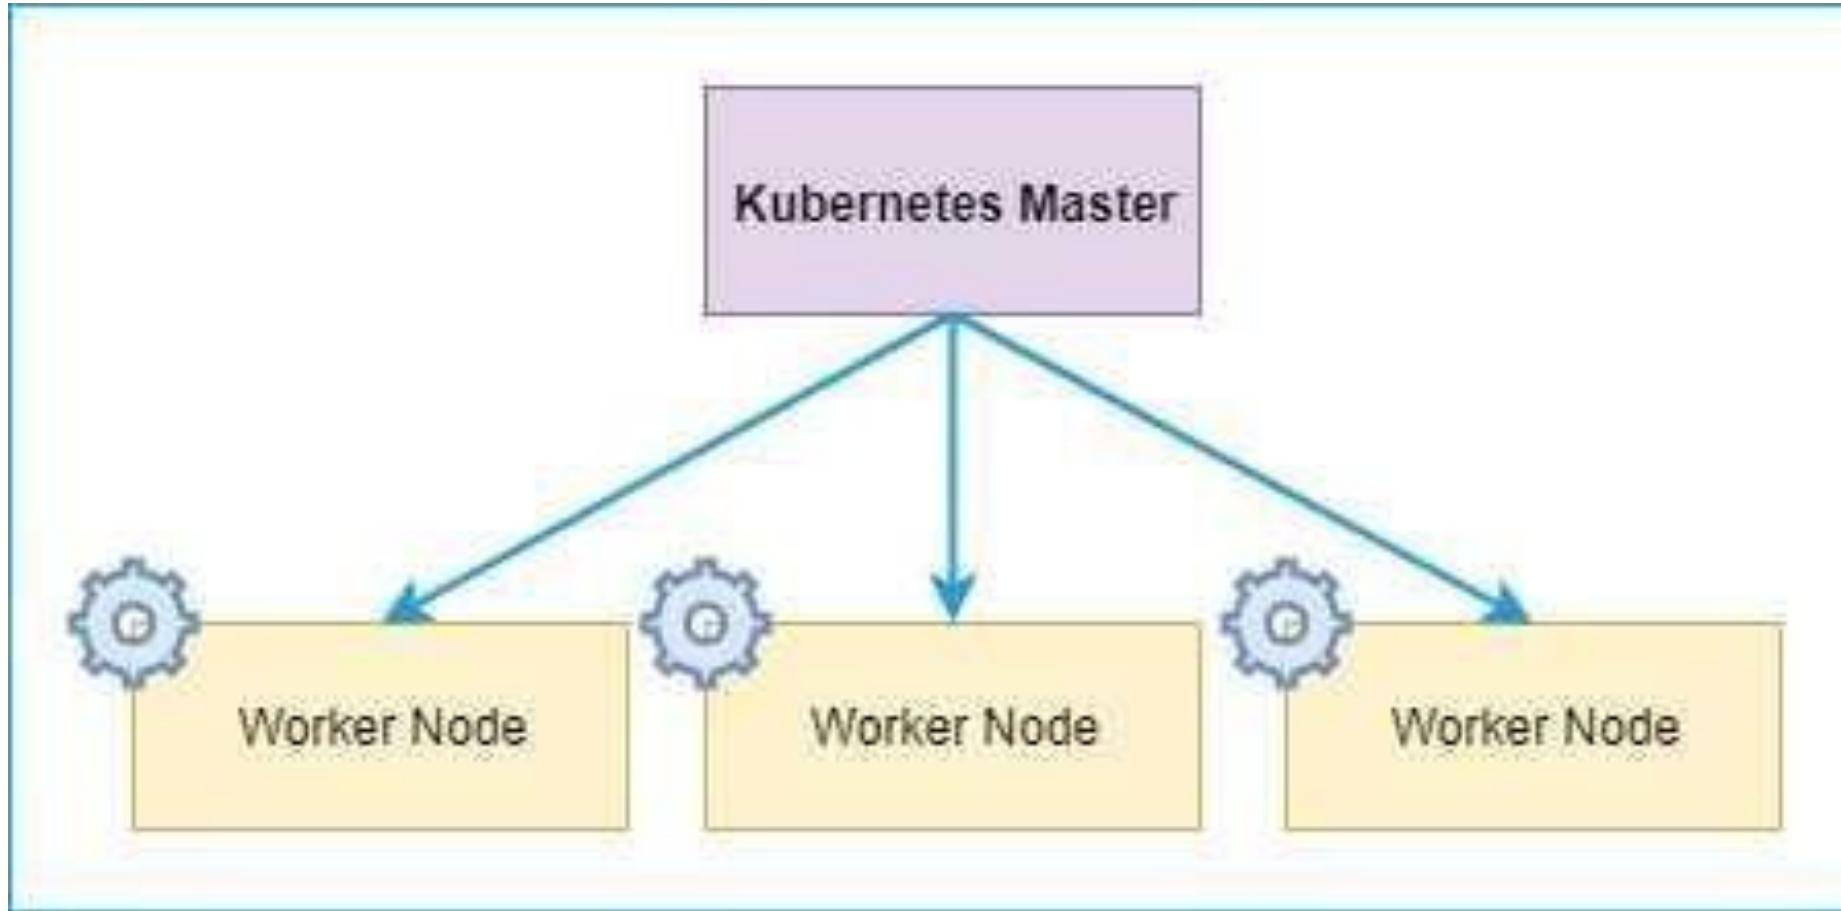
Auto Scaling



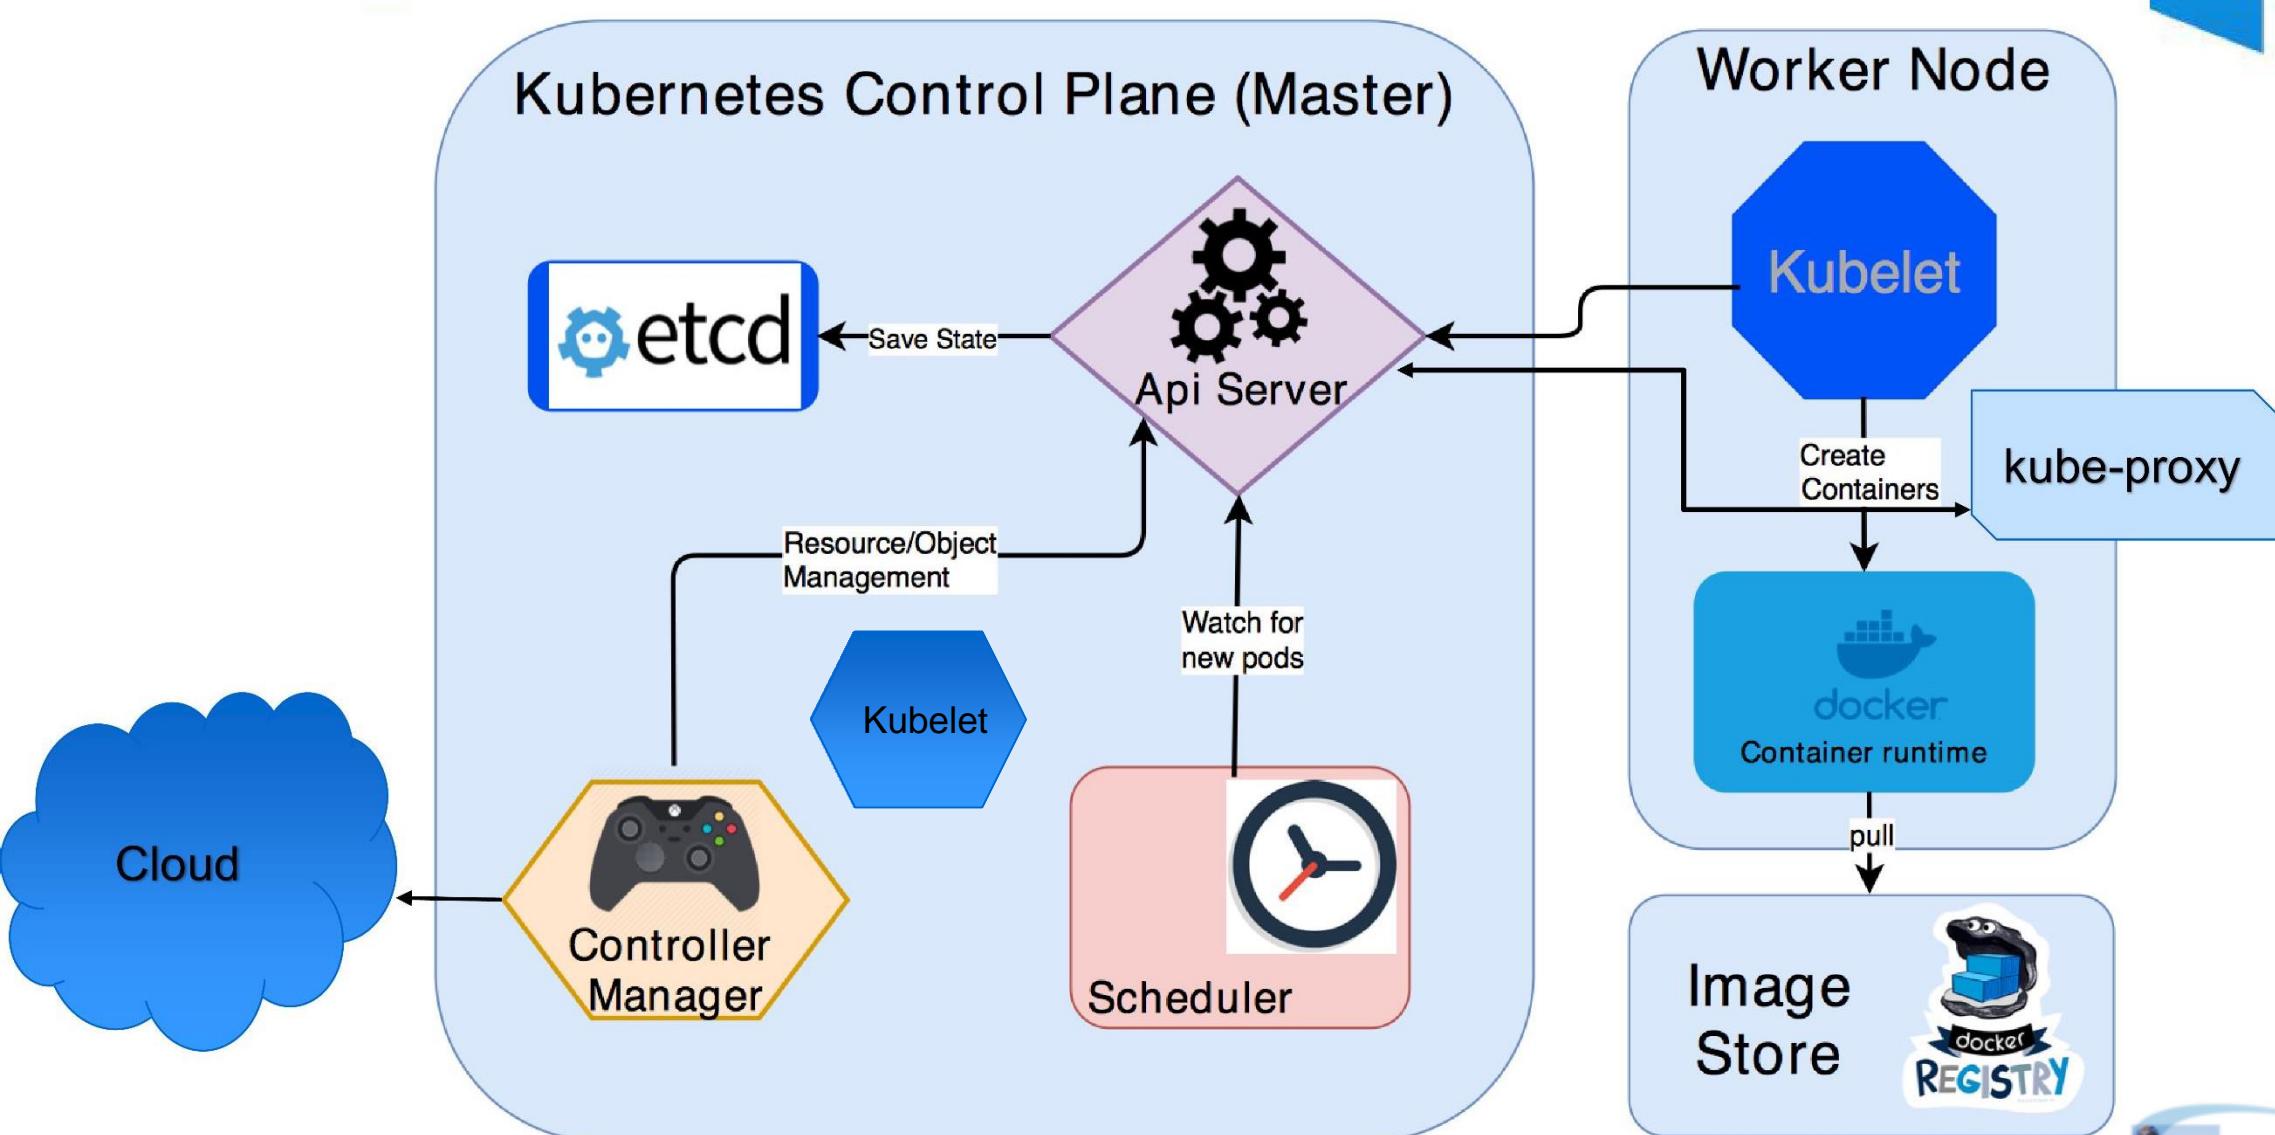
Load Balancing



# KUBERNETES COMPONENTS



# KUBERNETES COMPONENTS / ARCHITECTURE



# KUBE-APISERVER

Kubernetes is an API server which provides all the operation on cluster using the API. API server implements an interface, which means different tools and libraries can readily communicate with it. Kubeconfig is a package along with the server side tools that can be used for communication. It exposes Kubernetes API.

# I Kube-api Server

1. Authenticate User
2. Validate Request
3. Retrieve data
4. Update ETCD
5. Scheduler
6. Kubelet

# ETCD

- ETCD is a distributed reliable key-value store that is simple, secure and fast.
- It stores the configuration information which can be used by each of the nodes in the cluster. It is a high availability key value store that can be distributed among multiple nodes.
- It is storing all the information of master and worker nodes
- It is accessible only by Kubernetes API server as it may have some sensitive information.
- It is a distributed key value Store which is accessible to all.

# ETCD STORE information of

- 1) Nodes
- 2) POD's
- 3) Configs
- 4) Secrets
- 5) Accounts
- 6) Roles
- 7) Binding

# KUBE-CONTROLLER-MANAGER

Controllers are the brain behind the orchestration. They are responsible for noticing and responding when nodes, containers or end point goes down. The controllers make decision to make up new containers in such cases. This component is responsible for most of the collectors that regulates the state of cluster and performs a task.

# KUBE -CONTROLLER-MANAGER

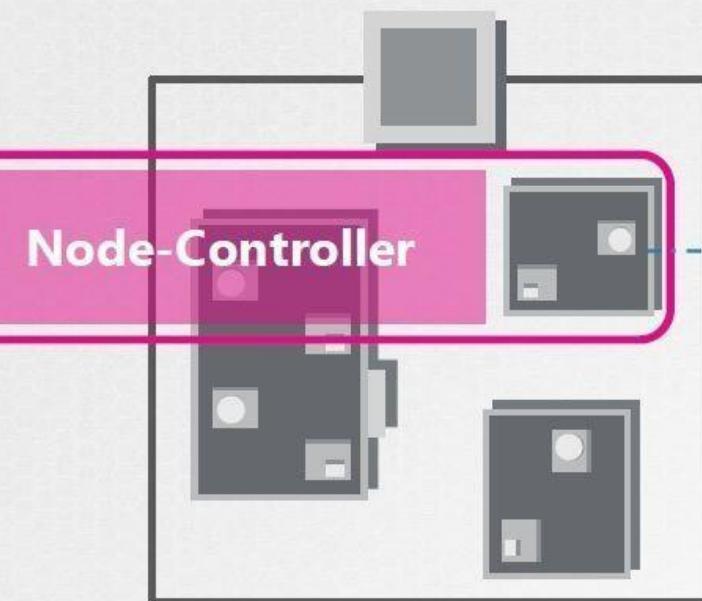
Some of the controllers that controller-manager include:

- **Node Controller:** Responsible for noticing and responding when nodes go down.
- **Replication Controller:** Responsible for maintaining the correct number of pods for every replication controller object in the system.
- **Endpoints Controller:** Populates the Endpoints object (that is, joins Services & Pods).
- **Service Account & Token Controllers:** Create default accounts and API access tokens for new namespaces.

# NODE CONTROLLER

- The node controller has multiple roles in a node's life. The first is assigning a CIDR block to the node when it is registered (if CIDR assignment is turned on).
- The second is keeping the node controller's internal list of nodes up to date with the cloud provider's list of available machines. When running in a cloud environment, whenever a node is unhealthy, the node controller asks the cloud provider if the VM for that node is still available. If not, the node controller deletes the node from its list of nodes.
- The third is monitoring the nodes' health. The node controller is responsible for updating the NodeReady condition of NodeStatus to ConditionUnknown when a node becomes unreachable

# Controller



▶ `kubectl get nodes`

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0



Watch Status

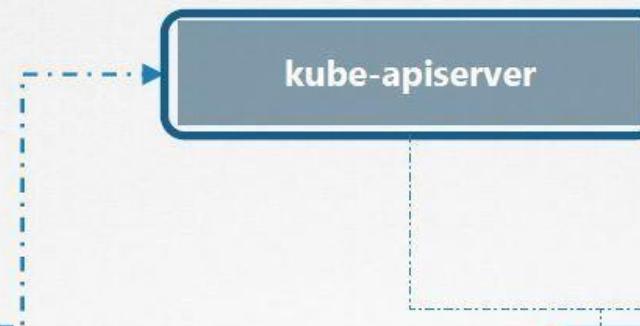
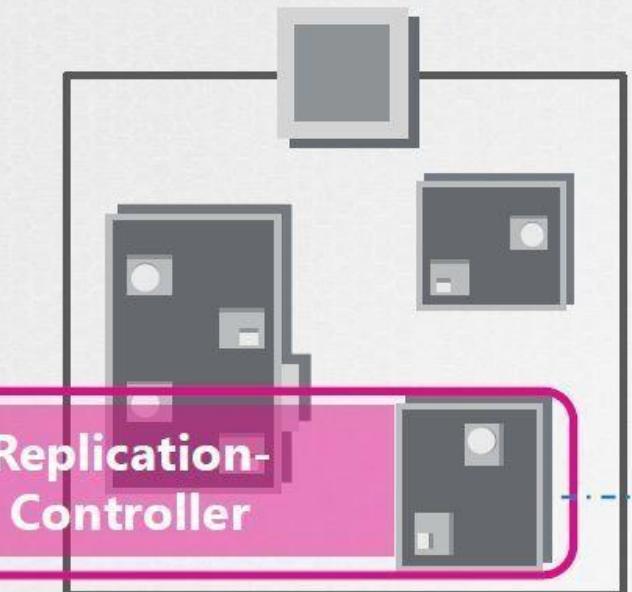
Remediate Situation

Node Monitor Period = 5s

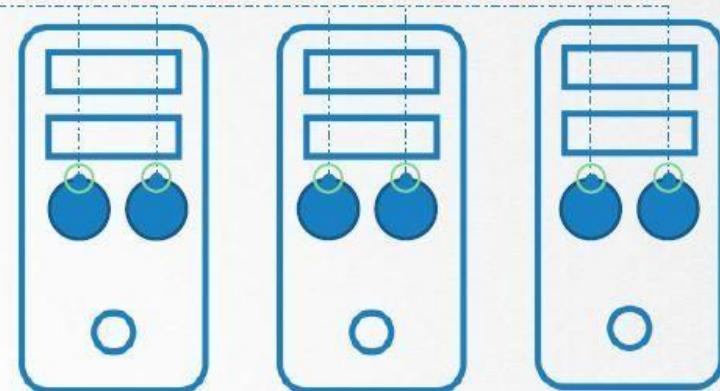
Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

# Controller



POD



Watch Status

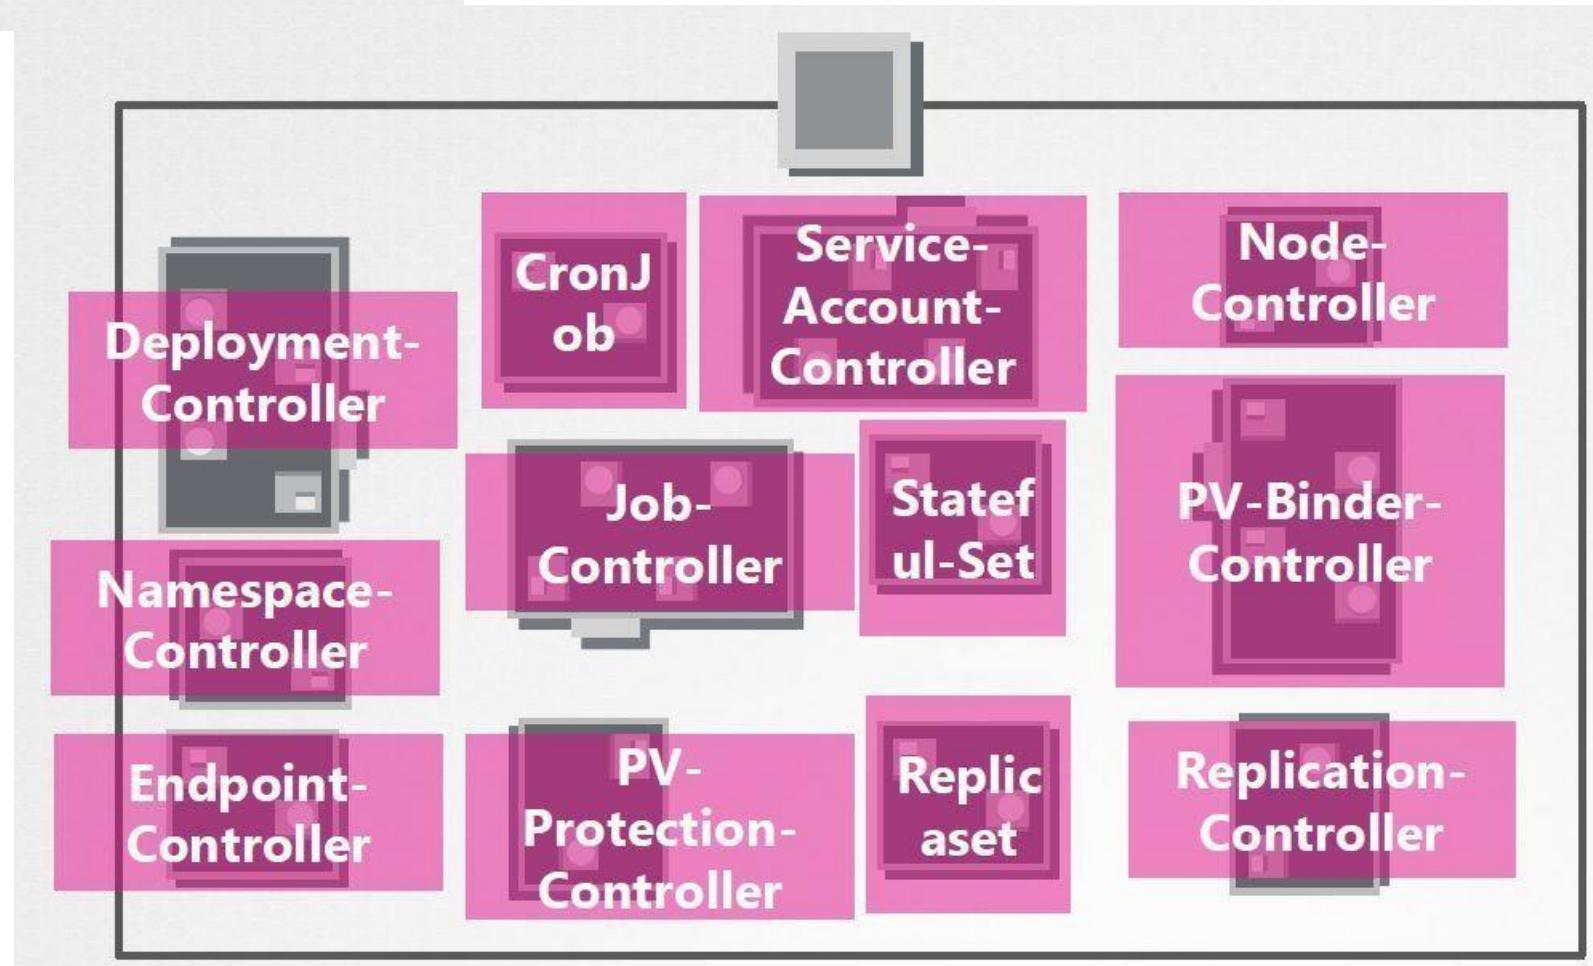
Remediate Situation

Node Monitor Period = 5s

Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

# I Controller



# KUBE-SCHEDULER

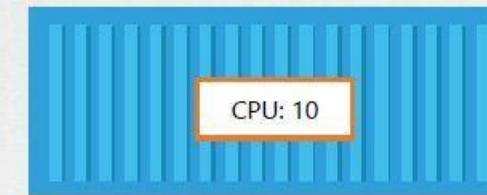
This is one of the key components of Kubernetes master. It is a service in master responsible for distributing the workload. It is responsible for tracking utilization of working load on cluster nodes and then placing the workload on which resources are available and accept the workload. In other words, this is the mechanism responsible for allocating pods to available nodes. The scheduler is responsible for workload utilization and allocating pod to new node.

# Kube-Scheduler

1. Filter Nodes

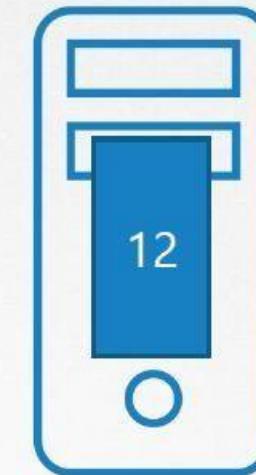


2. Rank Nodes



$$3^{\text{rd}} - 12 - 10 = 2$$

$$4^{\text{th}} - 16 - 10 = 6$$



# CONTAINER RUNTIME

- The container runtime is the software that is responsible for running containers.
- Kubernetes supports several container runtimes: Docker, containerd, CRI-O, and any implementation of the Kubernetes CRI (Container Runtime Interface).

# CONTAINER RUNTIME

containerd



cri-o

rkt



frakti

# KUBELET

It is the agent that runs on each node in the cluster. It is responsible for various tasks:

- making sure that the containers are running on the node as expected.
- for relaying information to and from control plane service through API server.
- It interacts with etcd store to read configuration details and wright values. This communicates with the master component to receive commands and work.
- The kubelet process then assumes responsibility for maintaining the state of work and the node server.
- It manages network rules, port forwarding, etc.



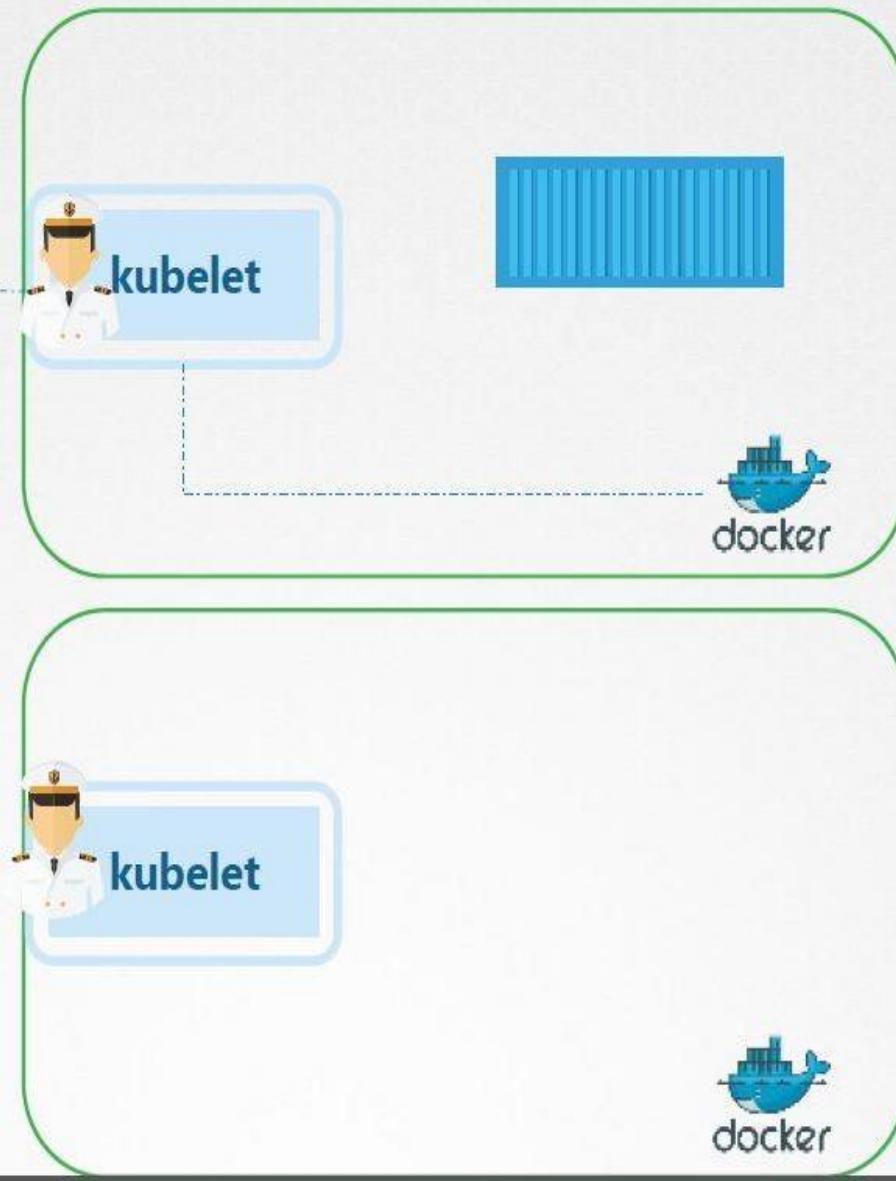
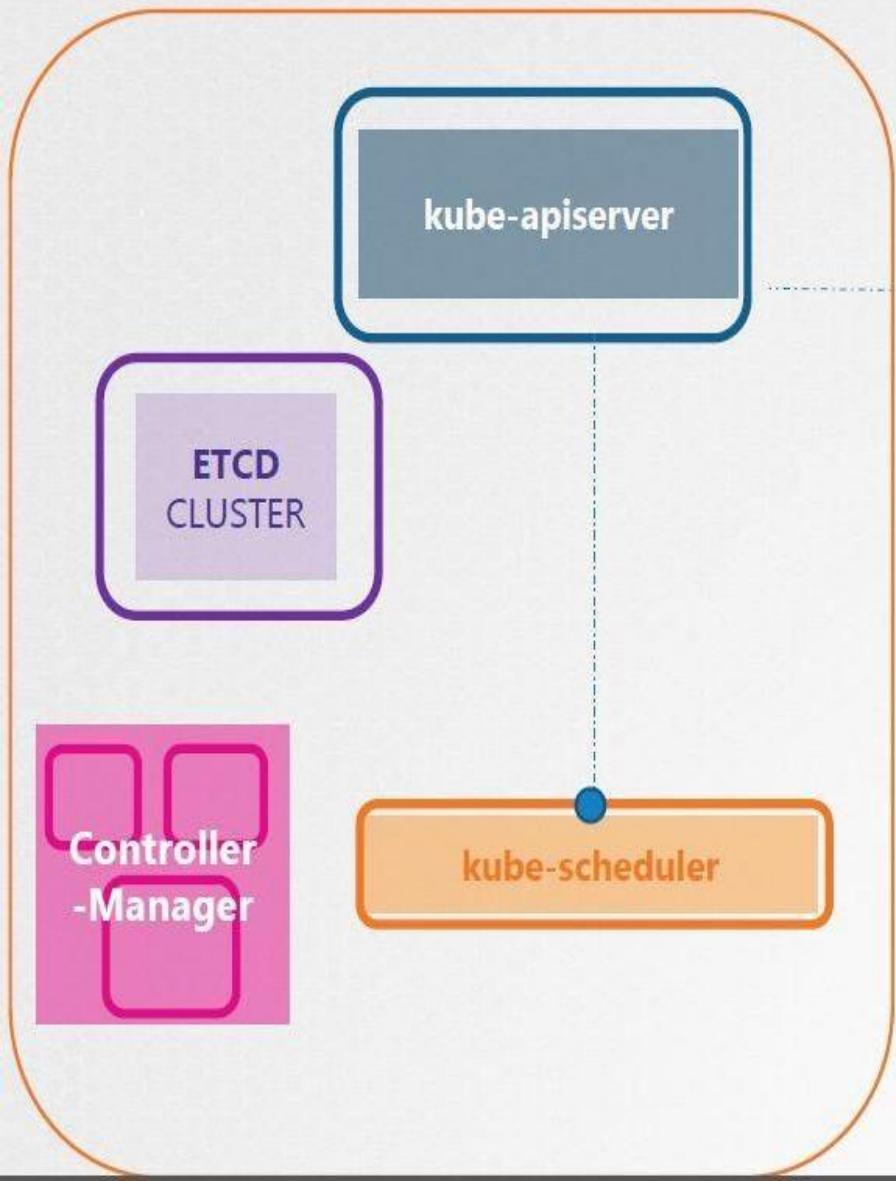
## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers



Register Node

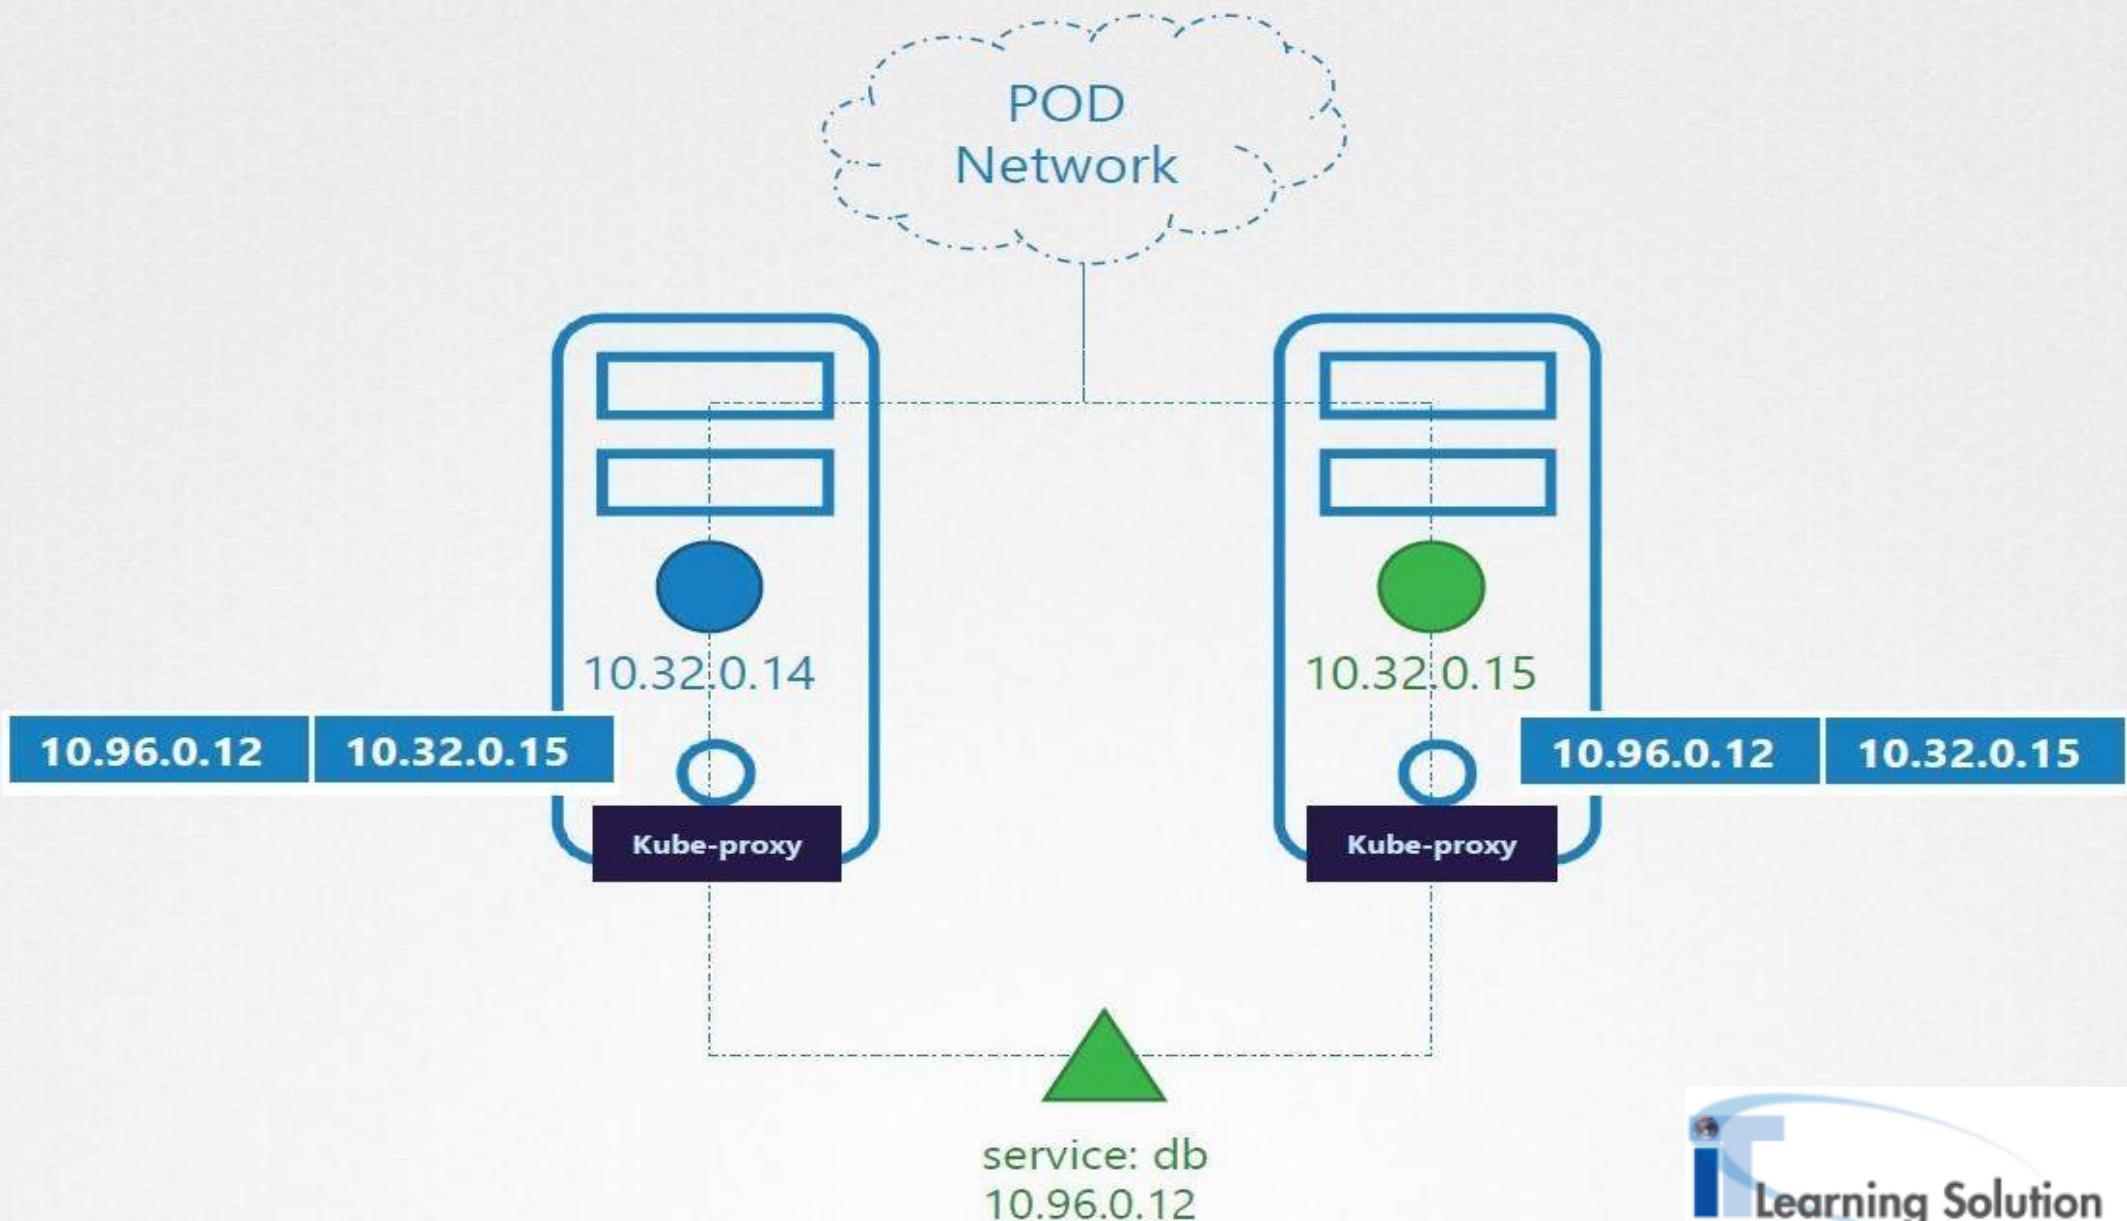
Create PODs

Monitor Node & PODs

# KUBE-PROXY

- This is a proxy service which runs on each node and helps in making services available to the external host. It helps in forwarding the request to correct containers and is capable of performing primitive load balancing. It makes sure that the networking environment is predictable and accessible and at the same time it is isolated as well. It manages pods on node, volumes, secrets, creating new containers' health checkup, etc.

# Kube-proxy



# KUBERNETES INSTALLATION

# REQUERIMENTES

One or more machines running one of:

- Ubuntu 16.04+
- Debian 9+
- CentOS 7
- Red Hat Enterprise Linux (RHEL) 7
- Fedora 25+
- HypriotOS v1.0.1+
- Container Linux (tested with 1800.6.0)

# REQUIREMENTS

- 2 GB or more of RAM per machine (any less will leave little room for your apps)
- 2 CPUs or more
- Full network connectivity between all machines in the cluster (public or private network is fine)
- Unique hostname, MAC address, and product\_uuid for every node.
- Certain ports are open on your machines.
- Swap disabled. You MUST disable swap in order for the kubelet to work properly.

# CHECK REQUIRED PORTS

## Control-plane node(s)

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	6443*	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10251	kube-scheduler	Self
TCP	Inbound	10252	kube-controller-manager	Self

## Worker node(s) ♂

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	30000-32767	NodePort Services†	All

# CONFIGURATION

## ON ALL NODES

- Disable SWAP

```
#vim /etc/fstab
```

then remove the entry of swap or comment

```
#swapoff -a -> to disable swap
```

- Set selinux to permissive mode

```
#vim /etc/sysconfig/selinux
```

```
SELINUX=permissive
```

```
#setenforce 0
```

- Stop firewall service

```
#systemctl stop firewalld
```

```
#systemctl disable firewalld
```

# INSTALL CONTAINER RUNTIME

## ON ALL NODES

- `#wget https://download.docker.com/linux/centos/docker-ce.repo -O /etc/yum.repos.d/docker-ce.repo`
- `# yum install docker-ce -y`
- `# systemctl start docker`
- `# systemctl enable docker`
- `# systemctl status docker`

# Installing kubeadm, kubelet and kubectl

- kubeadm: the command to bootstrap the cluster.
- kubelet: the component that runs on all of the machines in your cluster and does things like starting pods and containers.
- kubectl: the command line util to talk to your cluster.

# DEBIAN BASED

- `#apt-get update && sudo apt-get install -y apt-transport-https curl`
- `#curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -`
- `#cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list`  
`deb https://apt.kubernetes.io/ kubernetes-xenial main`  
EOF
- `#apt-get update`
- `#apt-get install -y kubelet kubeadm kubectl`
- `#apt-mark hold kubelet kubeadm kubectl`

# RPM BASED

## ON ALL NODES

```
#vim /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
```

```
#yum install kubelet kubectl kubeadm
#systemctl enable kubelet
#systemctl start kubelet
```

# NETWORK SOLUTION

We have many types of Network Solution->

- Calico
- Cilium
- Contiv-VPP
- Kube-router
- Weave Net

(doc -> <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/> )

# INITIALIZE THE CLUSTER

- Now initialize the master node->

```
#kubeadm init --pod-network-cidr=10.244.0.0/16
```

- Then after execution of above command you will have some commands so do all showing commands to configure kube configuration
- Run the join command in every worker node

# NETWORK SOLUTION

- we are using calico network for the installation

```
#kubectl apply -f https://docs.projectcalico.org/v3.14/manifests/calico.yaml
```

# VERIFY INSTALLATION

- `#kubectl get nodes`
- `#kubectl get pods`
- `#kubectl get pods -A`

# BASH COMPLETION

- `#yum install bash-completion -y`
- `#cd ~/.kube`
- `#kubectl completion bash > kube.sh`
- `#source kube.sh`

To make it permanent

- `#vim $HOME/.bashrc`
- Add a entry

`source $HOME/.kube/kube.sh`

# NODES

A node is a worker machine in Kubernetes, previously known as a minion. A node may be a VM or physical machine, depending on the cluster. Each node contains the services necessary to run pods and is managed by the master components.

# NODE STATUS

```
kubectl describe node <insert-node-name-here>
```

**A node's status contains the following information:**

- Addresses
- Conditions
- Capacity and Allocatable
- Info

# NODES STATE

Node Condition	Description
Ready	True if the node is healthy and ready to accept pods, False if the node is not healthy and is not accepting pods, and Unknown if the node controller has not heard from the node in the last node-monitor-grace-period (default is 40 seconds)
MemoryPressure	True if pressure exists on the node memory – that is, if the node memory is low; otherwise False
PIDPressure	True if pressure exists on the processes – that is, if there are too many processes on the node; otherwise False
DiskPressure	True if pressure exists on the disk size – that is, if the disk capacity is low; otherwise False
NetworkUnavailable	True if the network for the node is not correctly configured, otherwise False

# KUBECTL

```
#kubectl run <name> --image=<name>
#kubectl cluster-info
#kubectl get nodes
#kubectl describe node <insert-node-name-here>
#kubectl get pods
#kubectl describe pod <pod-name>
#kubectl get pods -A
#kubectl get pods -o wide
#kubectl get pods -A --show-labels
```

# PODS

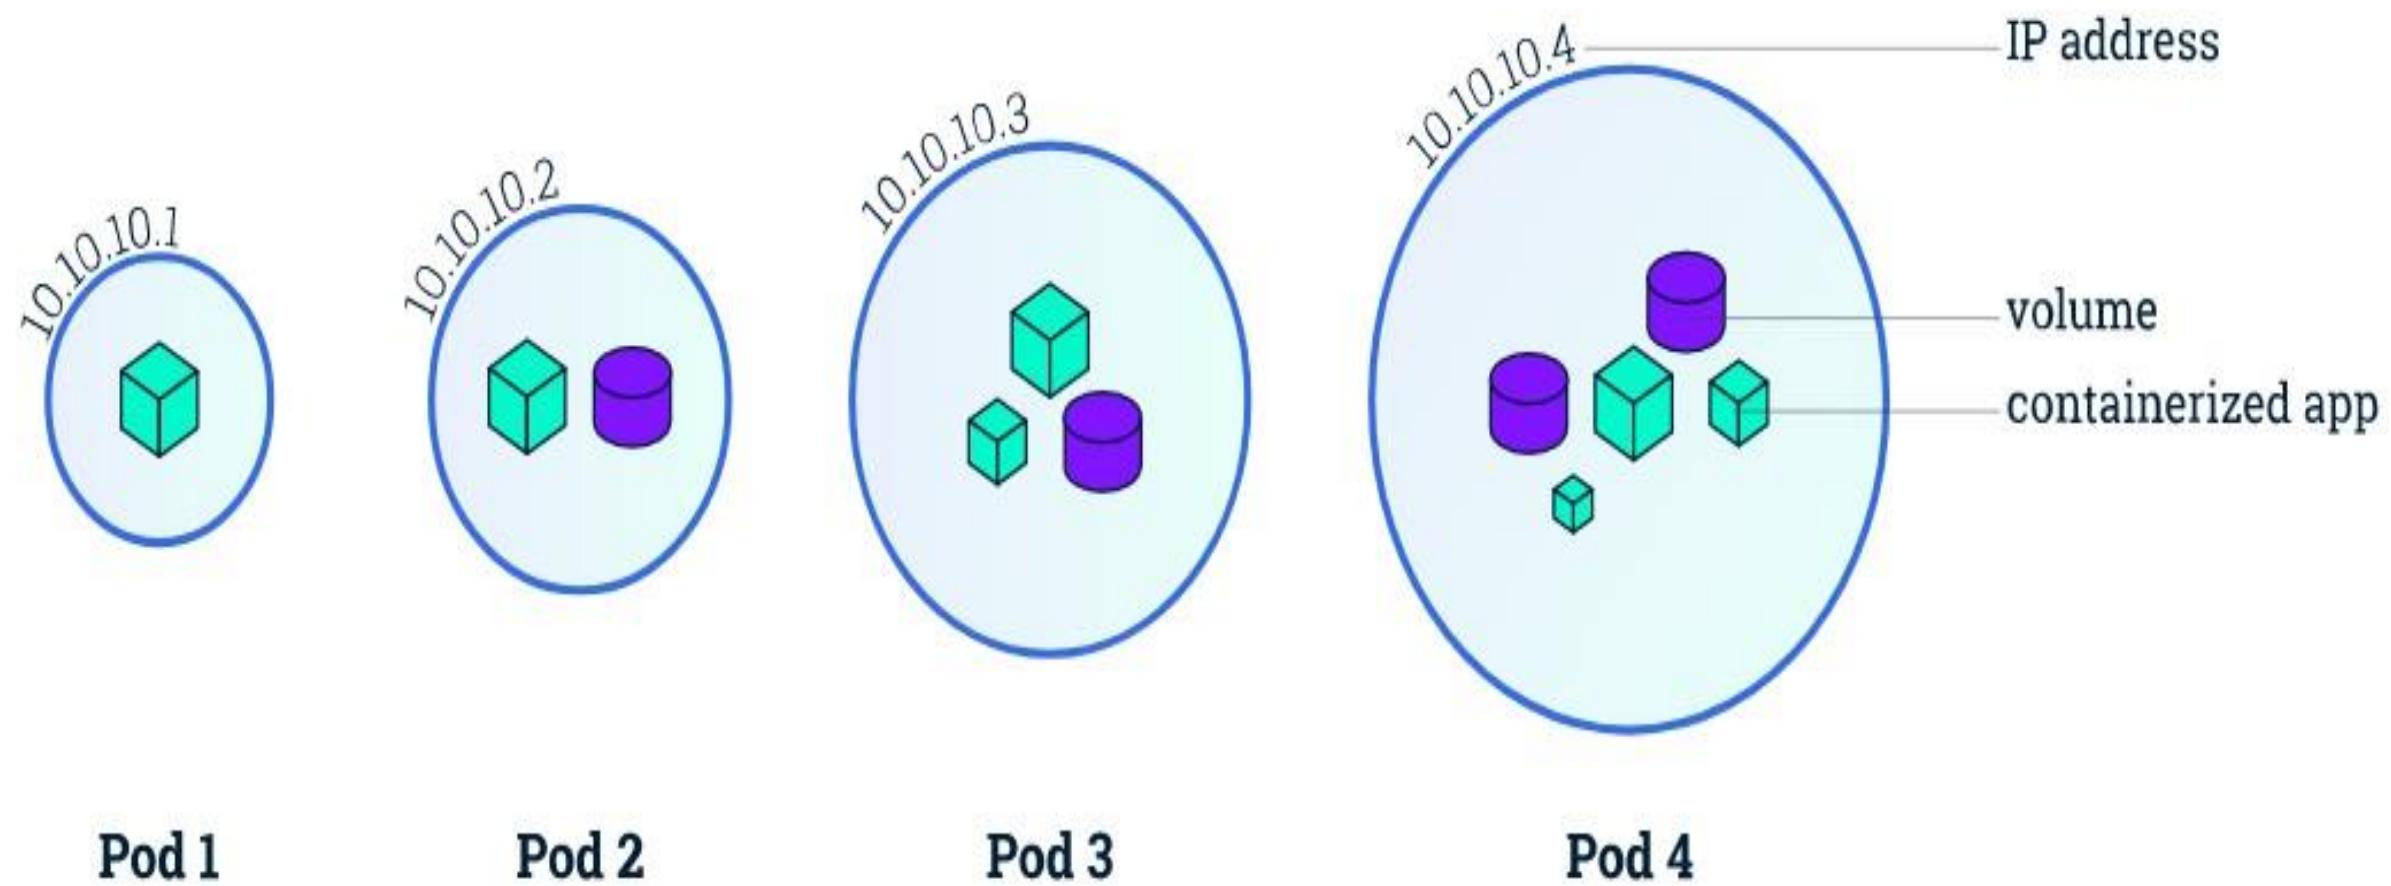
- A pod is a collection of containers and its storage inside a node of a Kubernetes cluster. It is possible to create a pod with multiple containers inside it. For example, keeping a database container and data container in the same pod.
- There are two types of Pods –
  - Single container pod
  - Multi container pod

# PODS

- Pods provide two kinds of shared resources for their constituent containers: networking and storage.

**1)Networking->** Each Pod is assigned a unique IP address. Every container in a Pod shares the network namespace, including the IP address and network ports.

**2)Storage ->** A Pod can specify a set of shared storage Volumes. All containers in the Pod can access the shared volumes, allowing those containers to share data.



# HOW TO CREATE PODS

- `#kubectl run <podname> --image=nginx`
- `#kubectl run <podname> --dry-run=client --image=nginx`
- `#kubectl get pods`
- `#kubectl get pods -A`
- `#kubectl get pods -o wide`
- `#kubectl get pods -A --show-labels`
- `#kubectl describe pods <podname>`
- `#kubectl explain <resource> like pod,deployment`
- `#kubectl exec -it <podname> bash`

# YAML INTRODUCTION

- Four fields are required to create yaml file->

**1)apiVersion** - Which version of the Kubernetes API you're using to create this object

**2) kind** - What kind of object you want to create

**3)metadata** - Data that helps uniquely identify the object, including a name string, UID, and optional namespace

**4) spec** - What state you desire for the object

# apiVersion

1) apiVersion: -> version of api depend on the module what we are creating at the time of creating yaml file we need to specify version (like:- service then in apiversion we will use v1 ) use like below options

<b>kind</b>	<b>Version</b>
Pod	v1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1

# HOW TO CREATE POD

```
#vim pod.yml
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: new-pod
5 spec:
6   containers:
7     - name: nginx-container
8       image: nginx
9
10
```

```
#kubectl create -f pod.yml
#kubectl create -f pod.yml --dry-run=client
#kubectl delete -f pod.yml
```

# INIT CONTAINER

- Init Containers are containers that run before the main container runs with your containerized application
- They are specialized containers that runs before app containers in a pod.
- Init containers can contain utilities or setup scripts not present in an app image.

Init containers are exactly like regular containers, except:

- Init containers always run to completion.
- Each init container must complete successfully before the next one starts.

**NOTE:- If a Pod's init container fails, Kubernetes repeatedly restarts the Pod until the init container succeeds.**

# INIT CONTAINER

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: firstcontainer
      image: nginx
      args: ["sleep","50"]
    - name: secondcontainer
      image: nginx
  initContainers:
    - name: initcontainer
      image: nginx
      args: ["sleep","30"]
```

# LABELS AND SELECTORS

## Labels :

- Labels enable end users to map their own, custom, organizational structures onto system resources in a loosely coupled fashion.

## Selectors :

- Labels do not provide uniqueness. In general, we can say many objects can carry the same labels. Labels selector are core grouping primitive in Kubernetes. They are used by the users to select a set of objects or to filter the tags.

# EXAMLE OF LABEL

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
  labels:
    env: development
spec:
  containers:
    - name: label-example
      image: nginx
      ports:
        - containerPort: 80
```

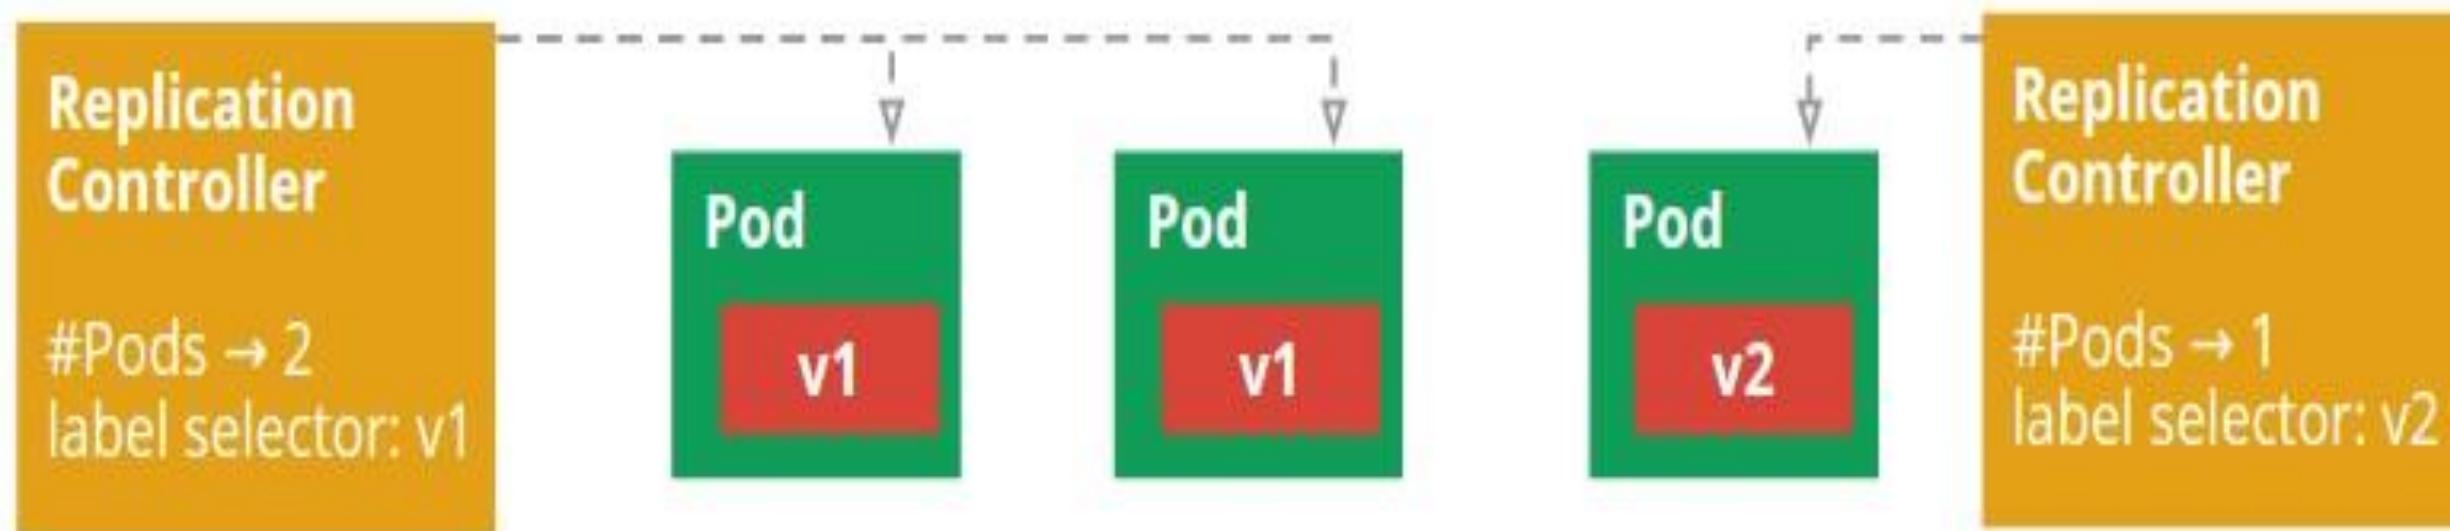
# LABELS COMMANDS

```
#kubectl apply -f <filename>
#kubectl get pods
#kubectl get pods <podname>--show-labels
#kubectl label pod <podname> tier=backend
#kubectl get pod <podname> --show-labels
#kubectl label pod <podname> tier-
#kubectl label --overwrite pods <podname> env=prod
```

# SELECTOR COMMANDS

- kubectl get pods -l env=prod
- kubectl get pods -l env!=prod

# Kubernetes Replication Controllers



## Behavior

- Keeps Pods running
- Gives direct control of Pod #s

## Benefits

- Restarts Pods, desired state
- Fine-grained control for scaling

# replica\_controller.yml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: rc-nginx
spec:
  replicas: 3
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
```

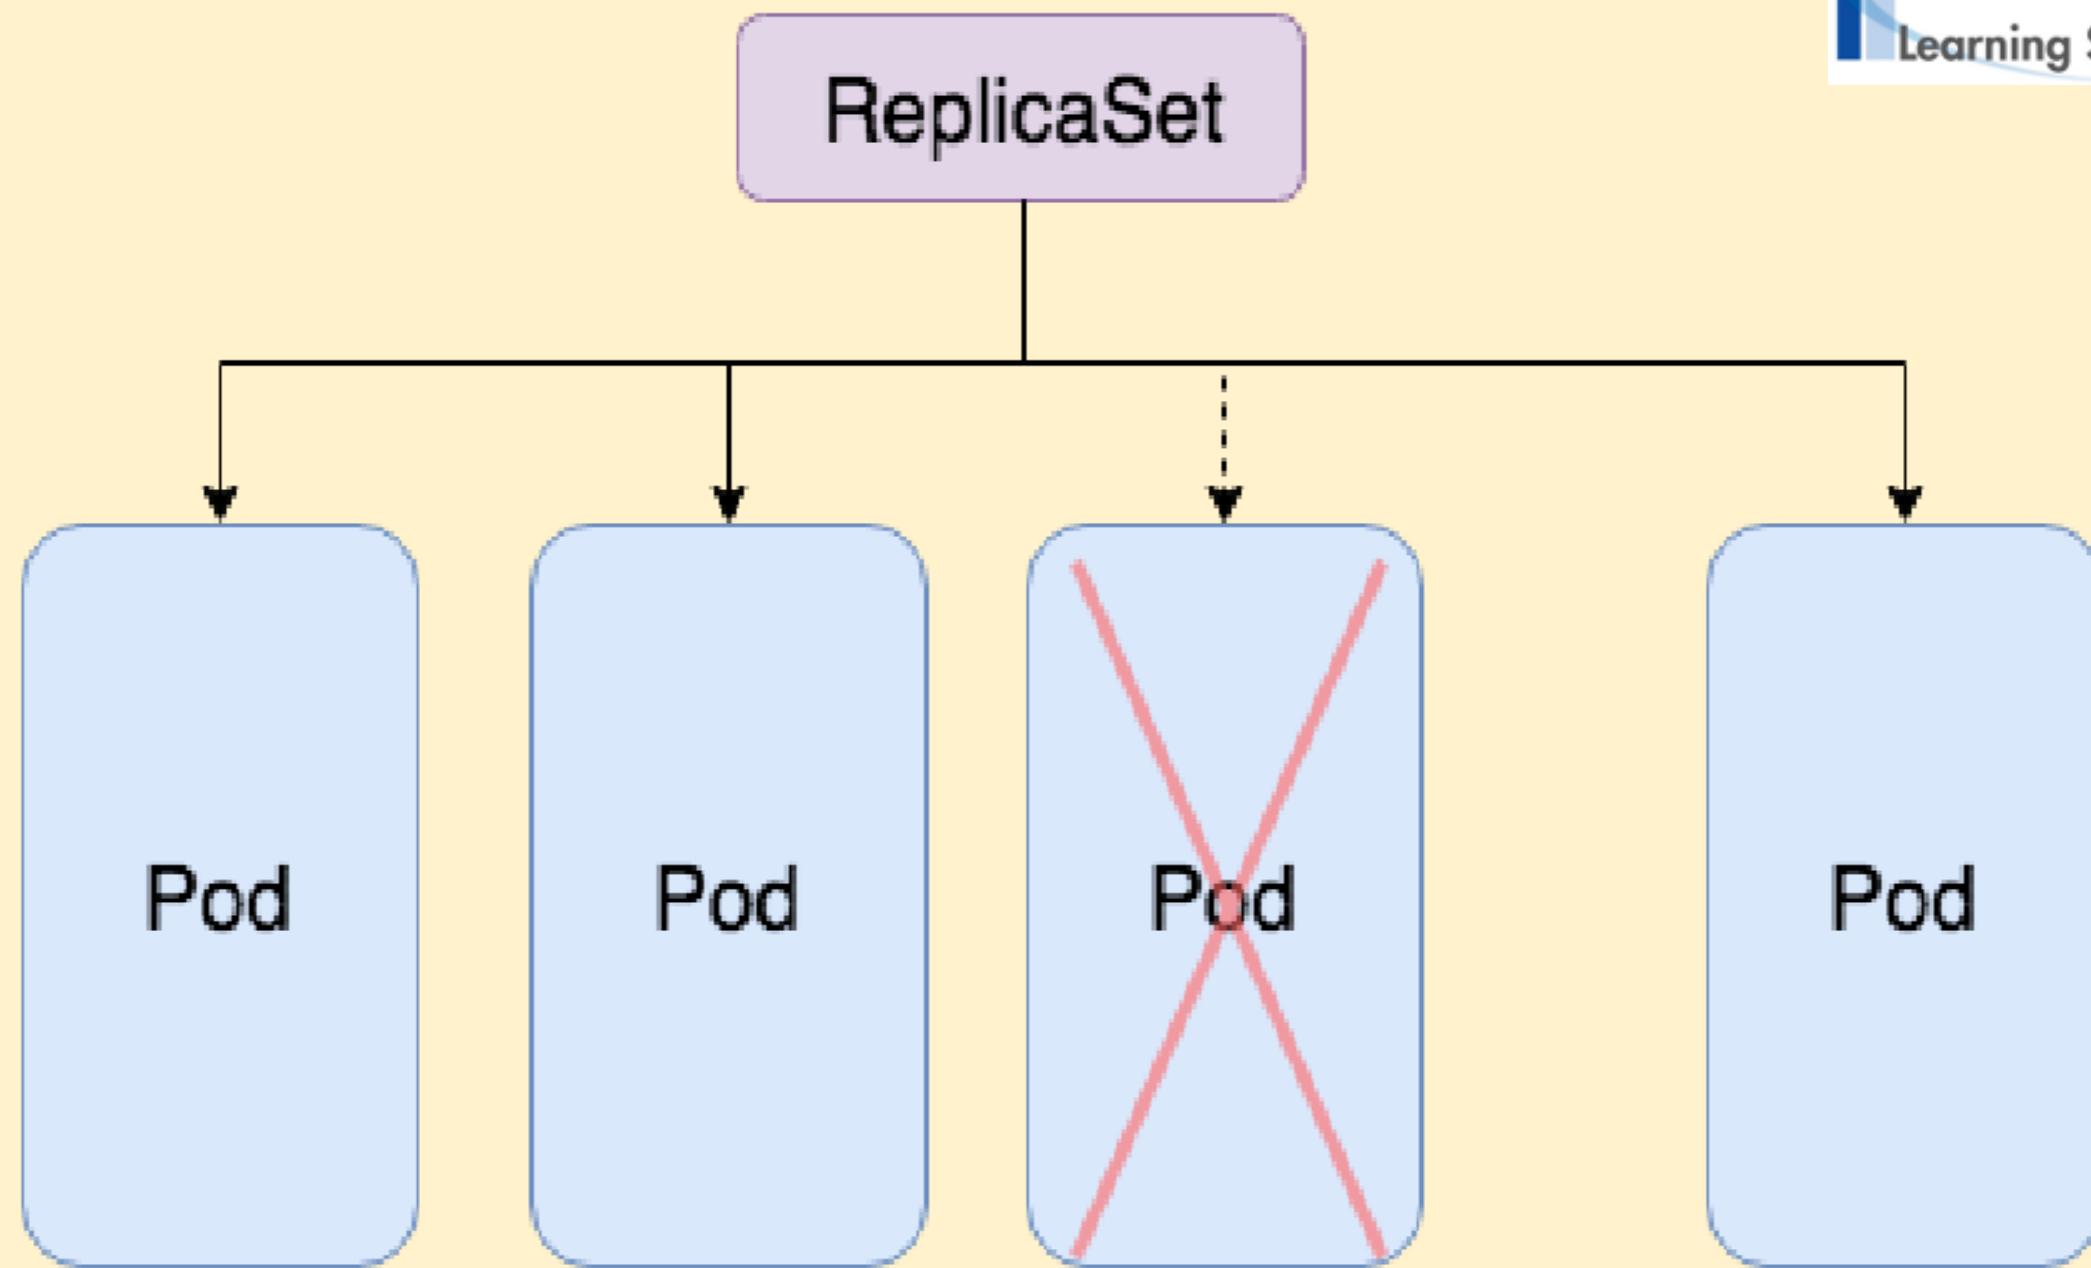
Pod definition

```
#kubectl create -f replica_controller.yml
#kubectl get replicationcontroller
or
#kubectl get rc
#kubectl get pods
#kubectl get pods -o wide
#kubectl get pods -A --show-labels
```

# ReplicaSet

- A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical Pods.
- How ReplicaSet Works?

A ReplicaSet is defined with fields, including a selector that specifies how to identify Pods it can acquire, a number of replicas indicating how many Pods it should be maintaining, and a pod template specifying the data of new Pods it should create to meet the number of replicas criteria. A ReplicaSet then fulfills its purpose by creating and deleting Pods as needed to reach the desired number. When a ReplicaSet needs to create new Pods, it uses its Pod template.



# replica\_set.yml

```
1 apiVersion: apps/v1
2 kind: ReplicaSet
3 metadata:
4   name: rs-nginx
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: nginx
10  template:
11    metadata:
12      name: nginx
13      labels:
14        app: nginx
15    spec:
16      containers:
17      - name: nginx
18        image: nginx
19        ports:
20        - containerPort: 80
```

} Will match the labels from already created pods also

#kubectl create -f replica\_set.yml

#kubectl get replicaset  
or

#kubectl get rs  
#kubectl get pods  
#kubectl get pods -o wide  
#kubectl get pods -A --show-labels

If we want to increase or decrease the replicas then just modify the definition file  
replicas: 6

#kubectl replace -f replica\_set.yml

or

#kubectl scale --replicas=6 -f replica\_set.yml

or

#kubectl scale --replicas=6 replicaset rs-nginx

# REPLICATION CONTROLLER VS REPLICA SET

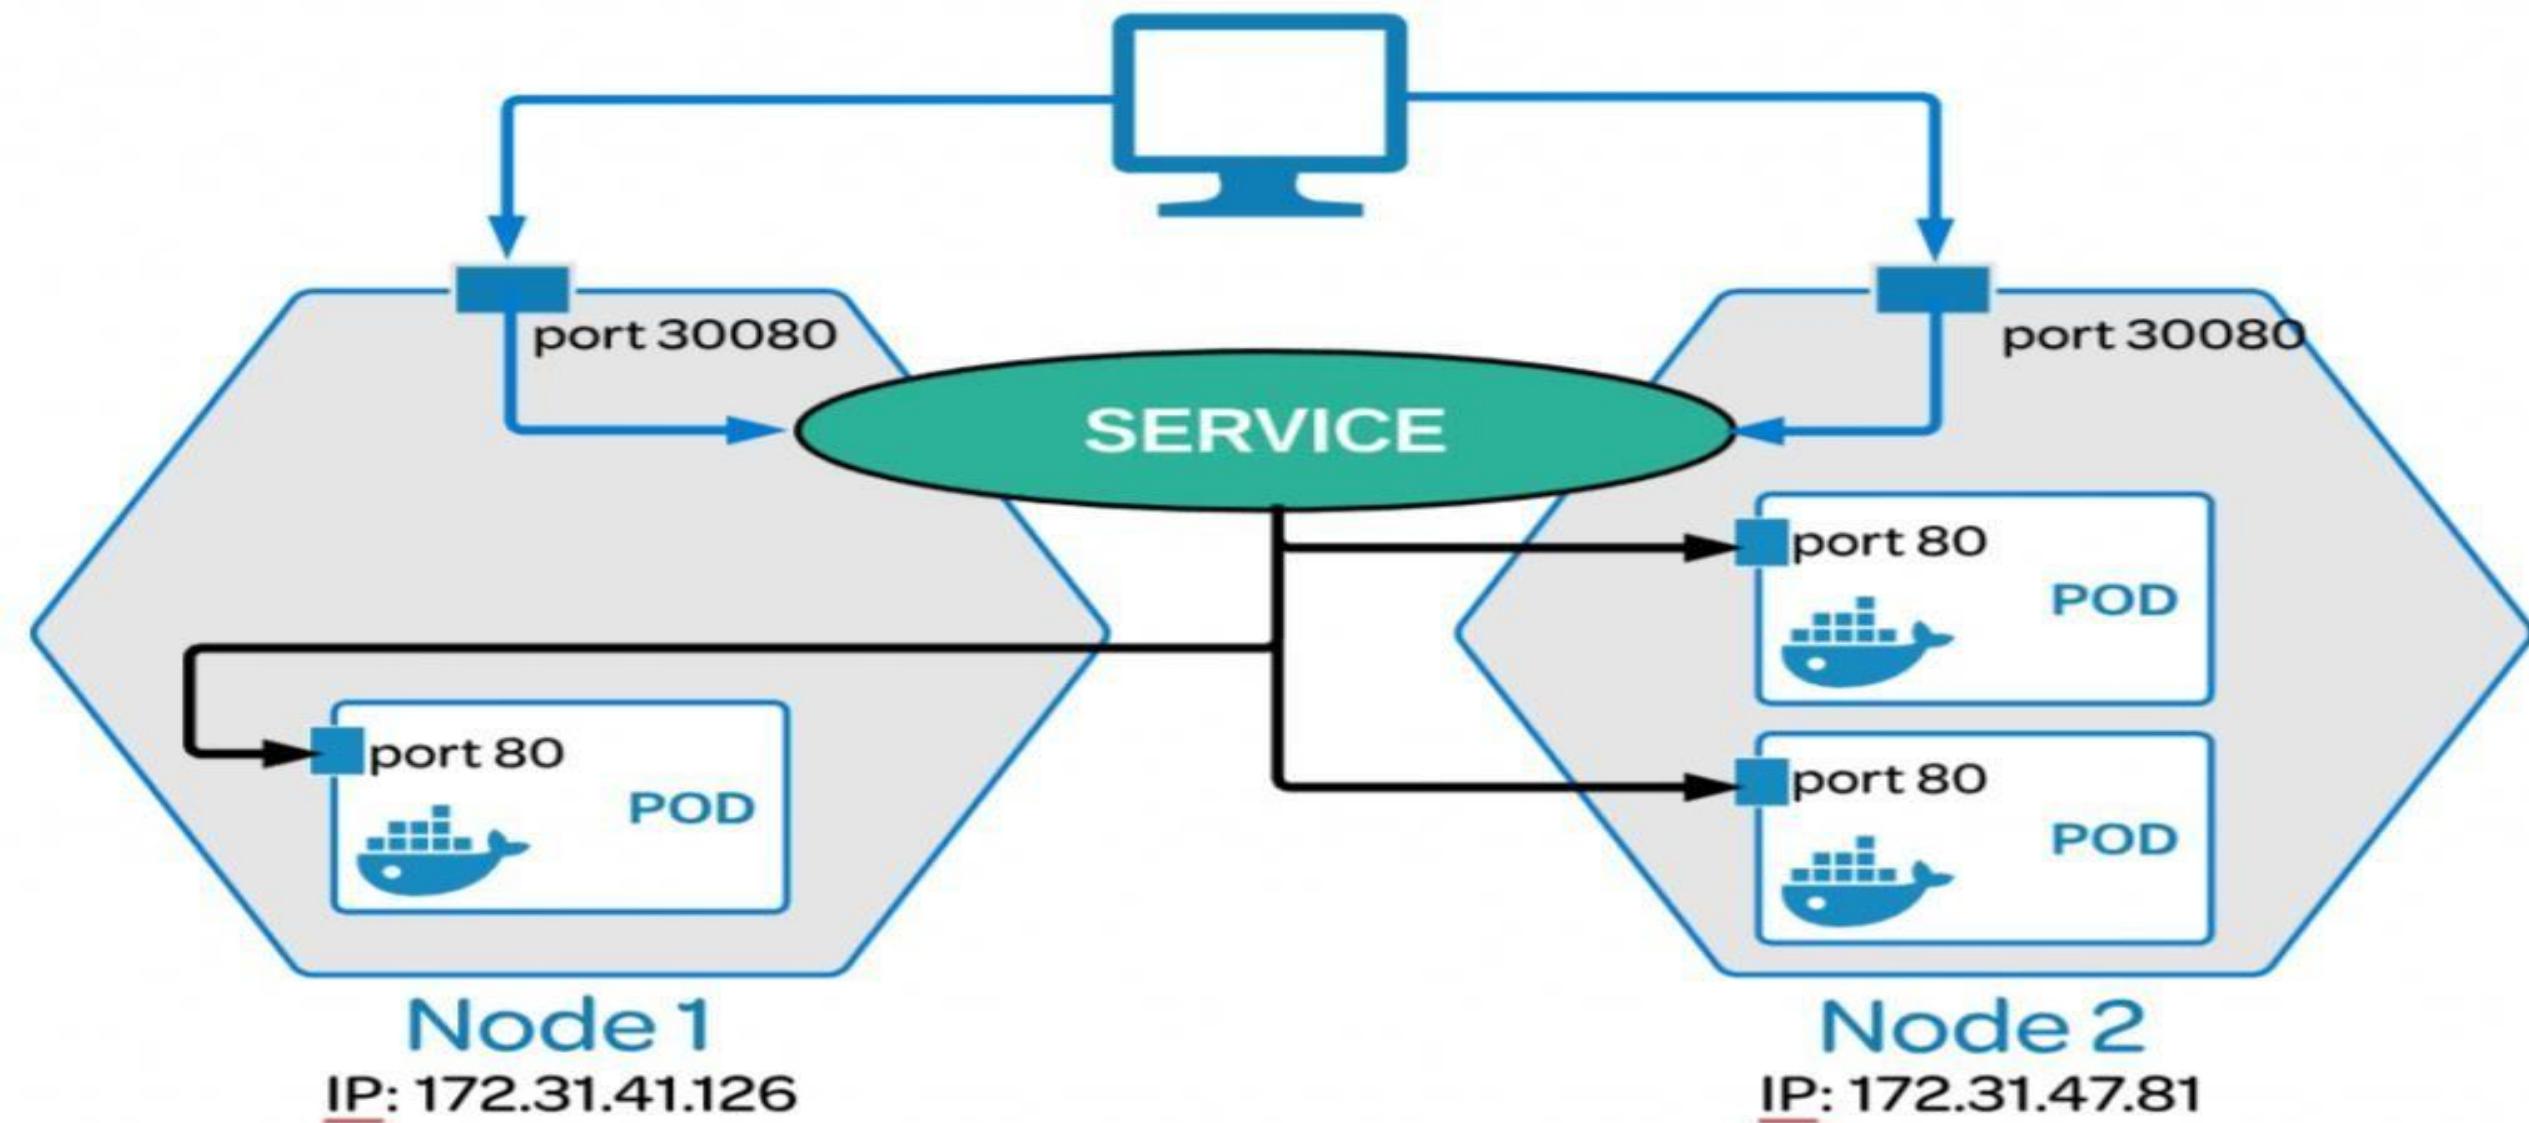
- ReplicaSet requires a selector definition. The selector helps to identify what pods are already created.
- Selector definition is not required for replication controller but we can specify.
- For example the pods created before the creation of replicaset that match same labels specified in the selector the replica set will automatically also take those pods into consideration when creating the replica set.

# SERVICE

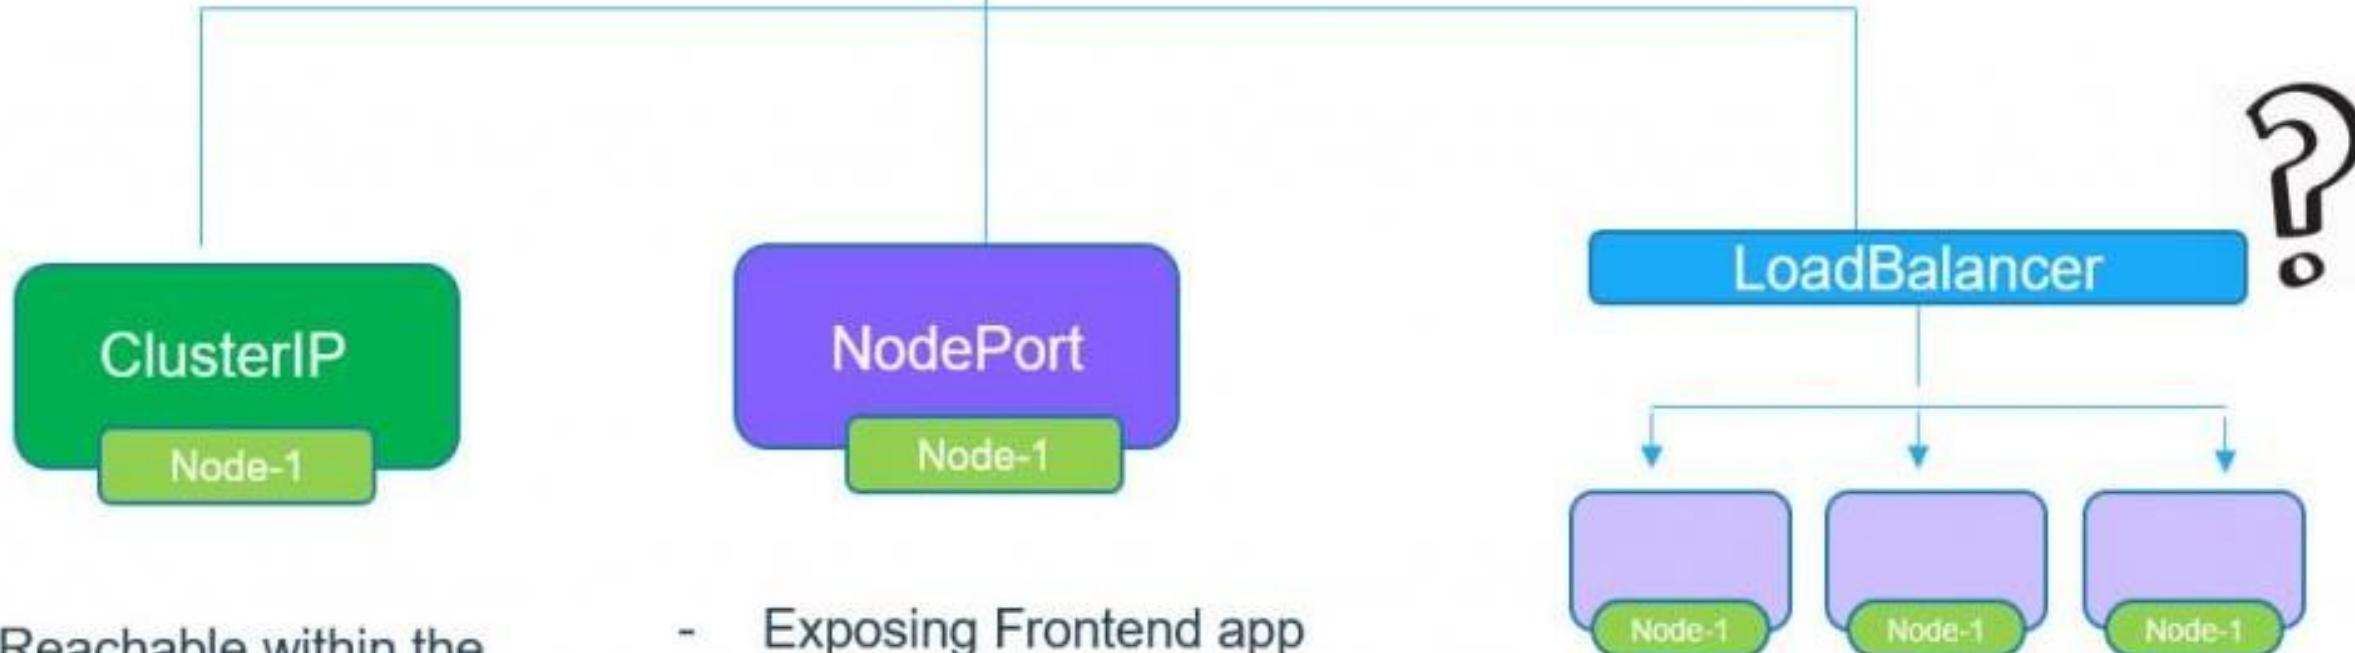
- Services are resources in Kubernetes (like pods, deployments, etc.) that provide a single point of access from the outside world, into your pod(s) which run your application.
- Each service has an IP address and a port that never change, and the client always knows where to find the pods due to the flat nat-less network.

# Kubernetes Service

A service allows you to dynamically access a group of replica pods.



## Types of Services

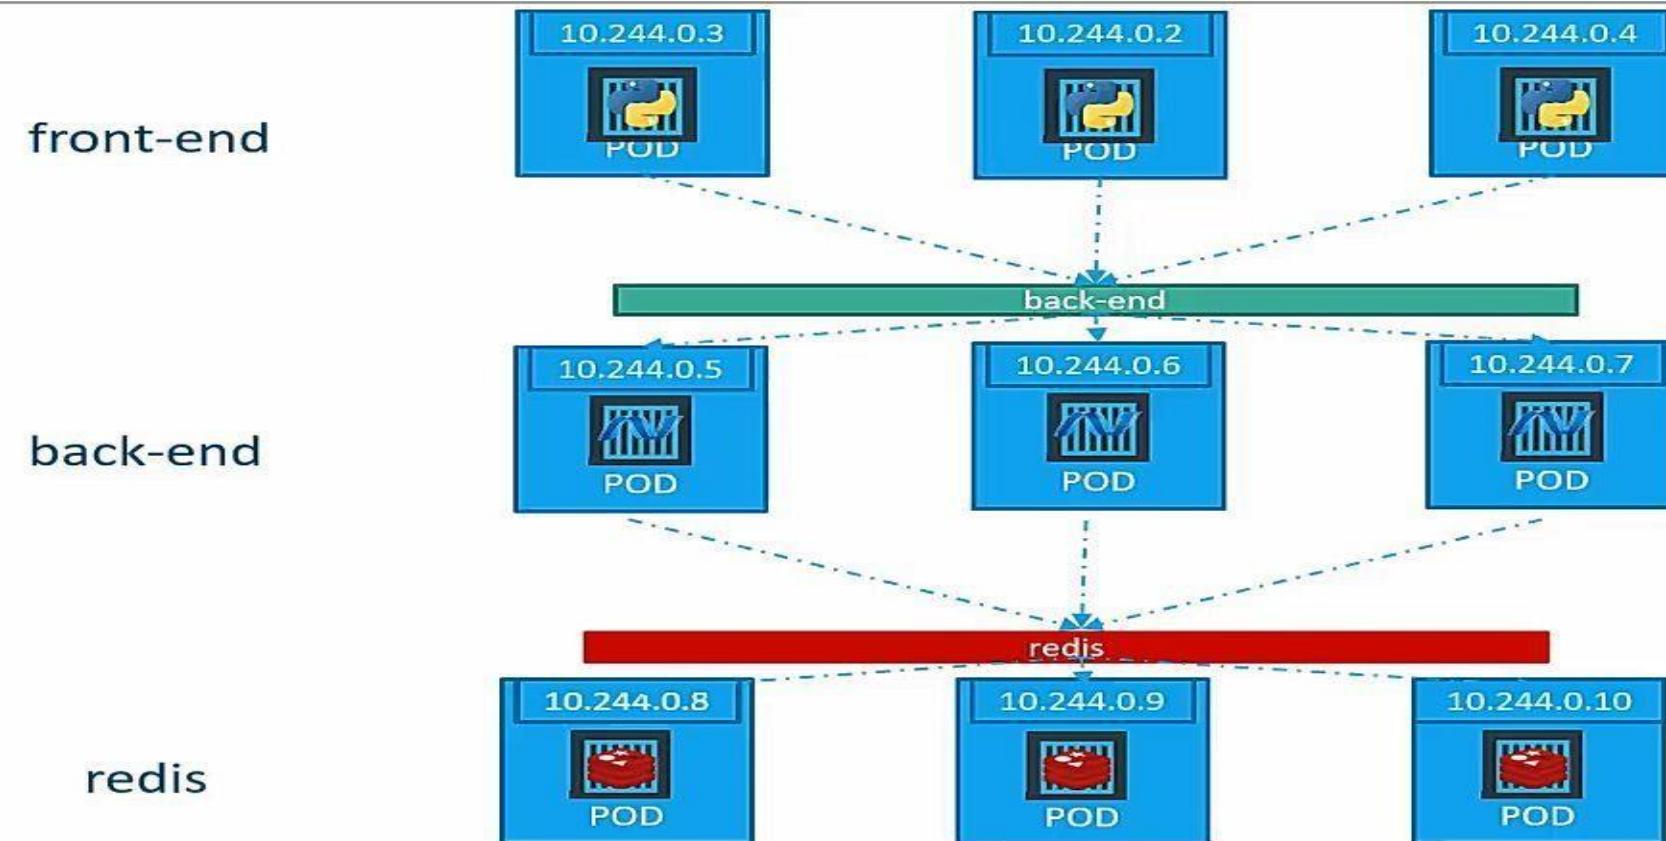


- Reachable within the cluster.
- Connects Frontend Pods to Backend Pods
- Exposing Frontend app to external world
- Equally distribute the loads

# CLUSTER IP

- ClusterIP: Exposes the Service on a cluster-internal IP. Choosing this value makes the Service only reachable from within the cluster. This is the default ServiceType.

# ClusterIP



# APP

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  labels:
    app: web
spec:
  containers:
    - name: web-container
      image: nginxdemos/hello:plain-text
      ports:
        - containerPort: 80
```

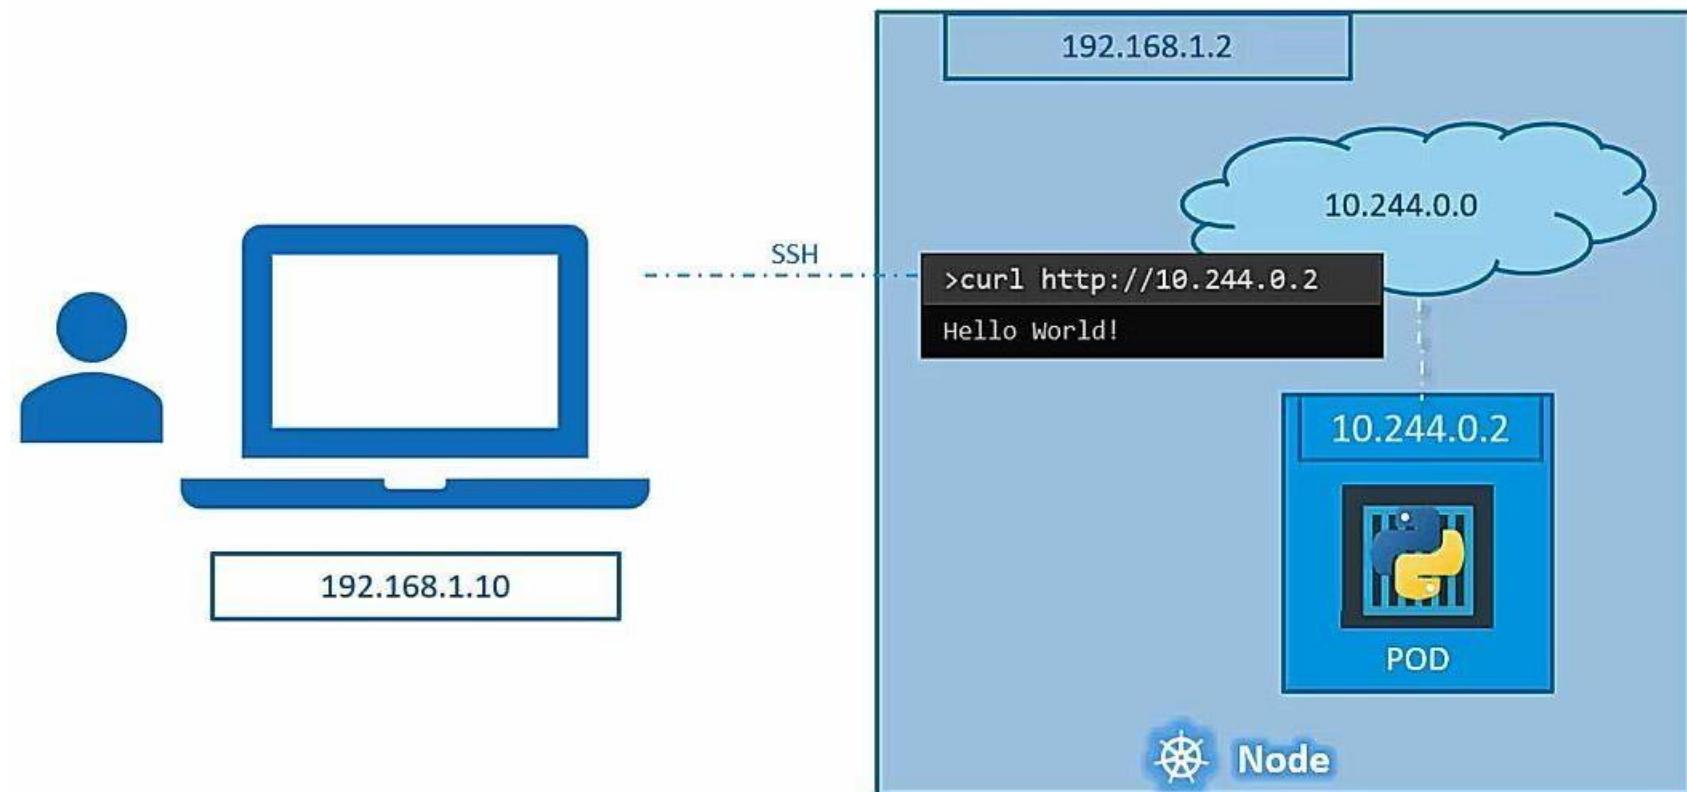
# CLUSTER IP SERVICE

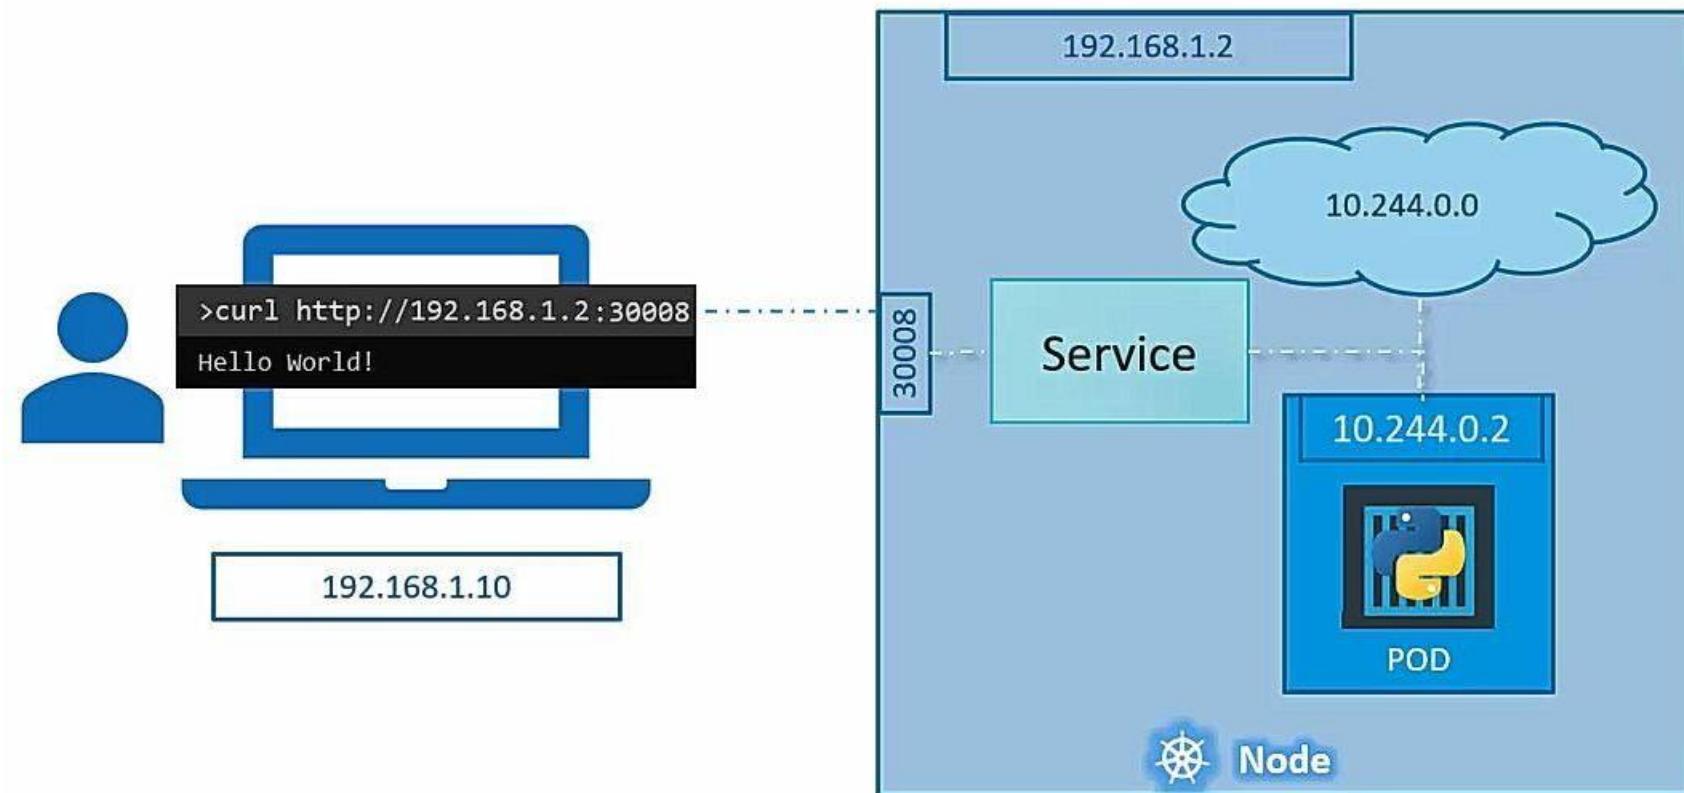
```
apiVersion: v1
kind: Service
metadata:
  name: my-pod-service
spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80

  selector:
    app: web
```

# NODE PORT

- NodePort: Exposes the Service on each Node's IP at a static port (the NodePort). A ClusterIP Service, to which the NodePort Service routes, is automatically created. You'll be able to contact the NodePort Service, from outside the cluster, by requesting <NodeIP>:<NodePort>.

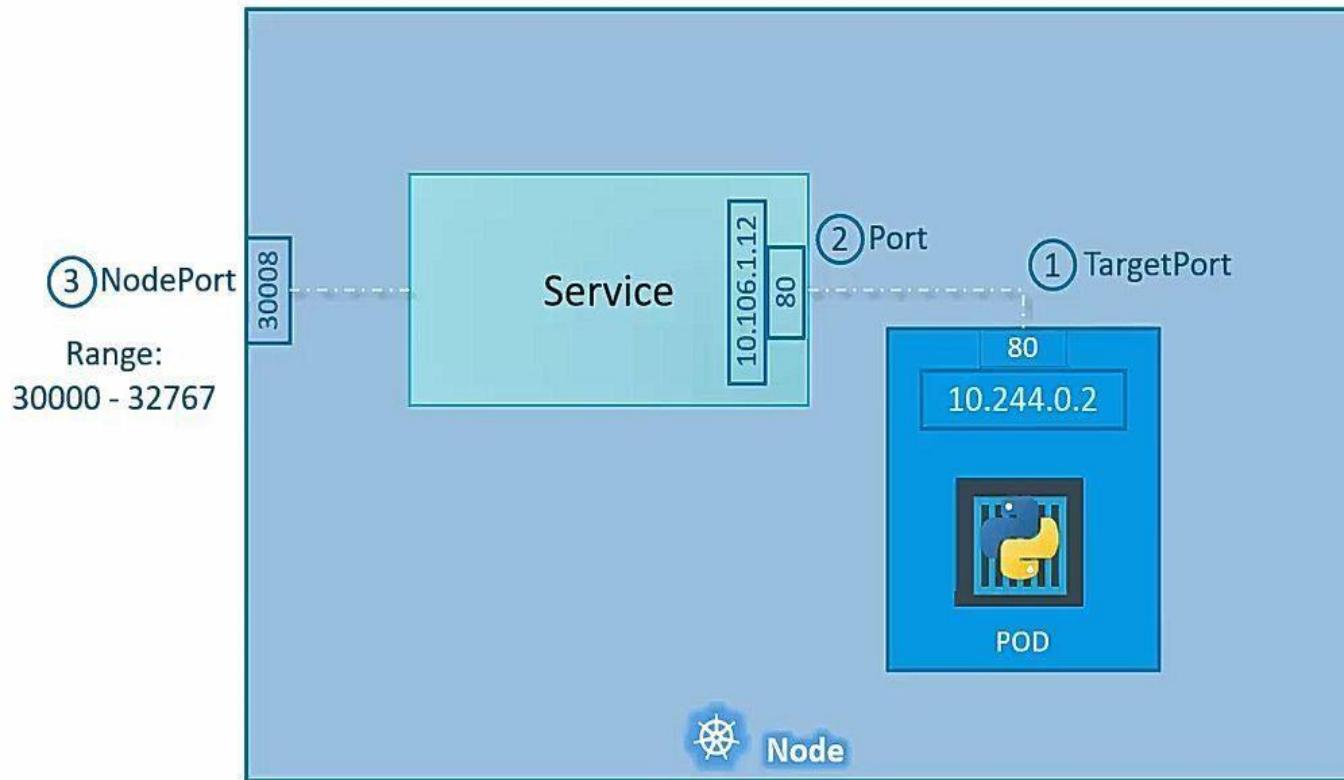




# APP

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  labels:
    app: web
spec:
  containers:
    - name: web-container
      image: nginxdemos/hello:plain-text
      ports:
        - containerPort: 80
```

# Service - NodePort



service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

# Service - NodePort

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: myapp
    type: front-end
```

```
> kubectl create -f service-definition.yml
service "myapp-service" created
```

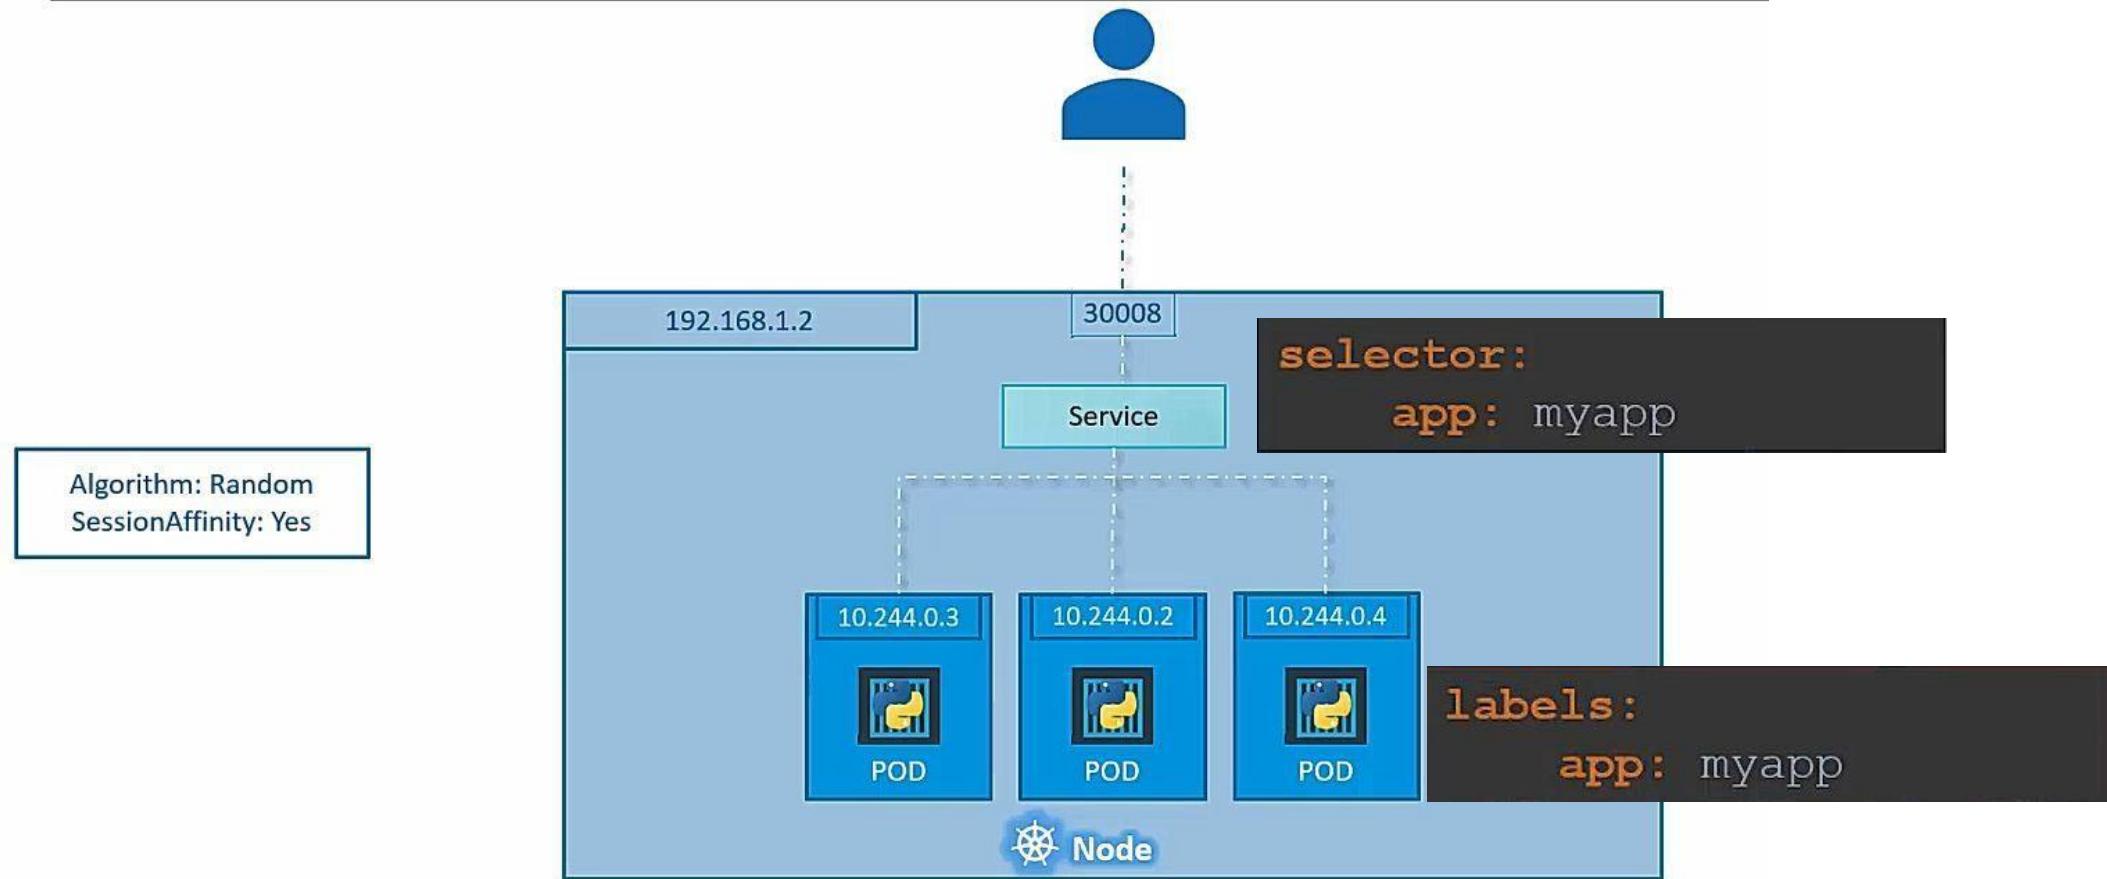
```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
myapp-service	NodePort	10.106.127.123	<none>	80:30008/TCP	5m

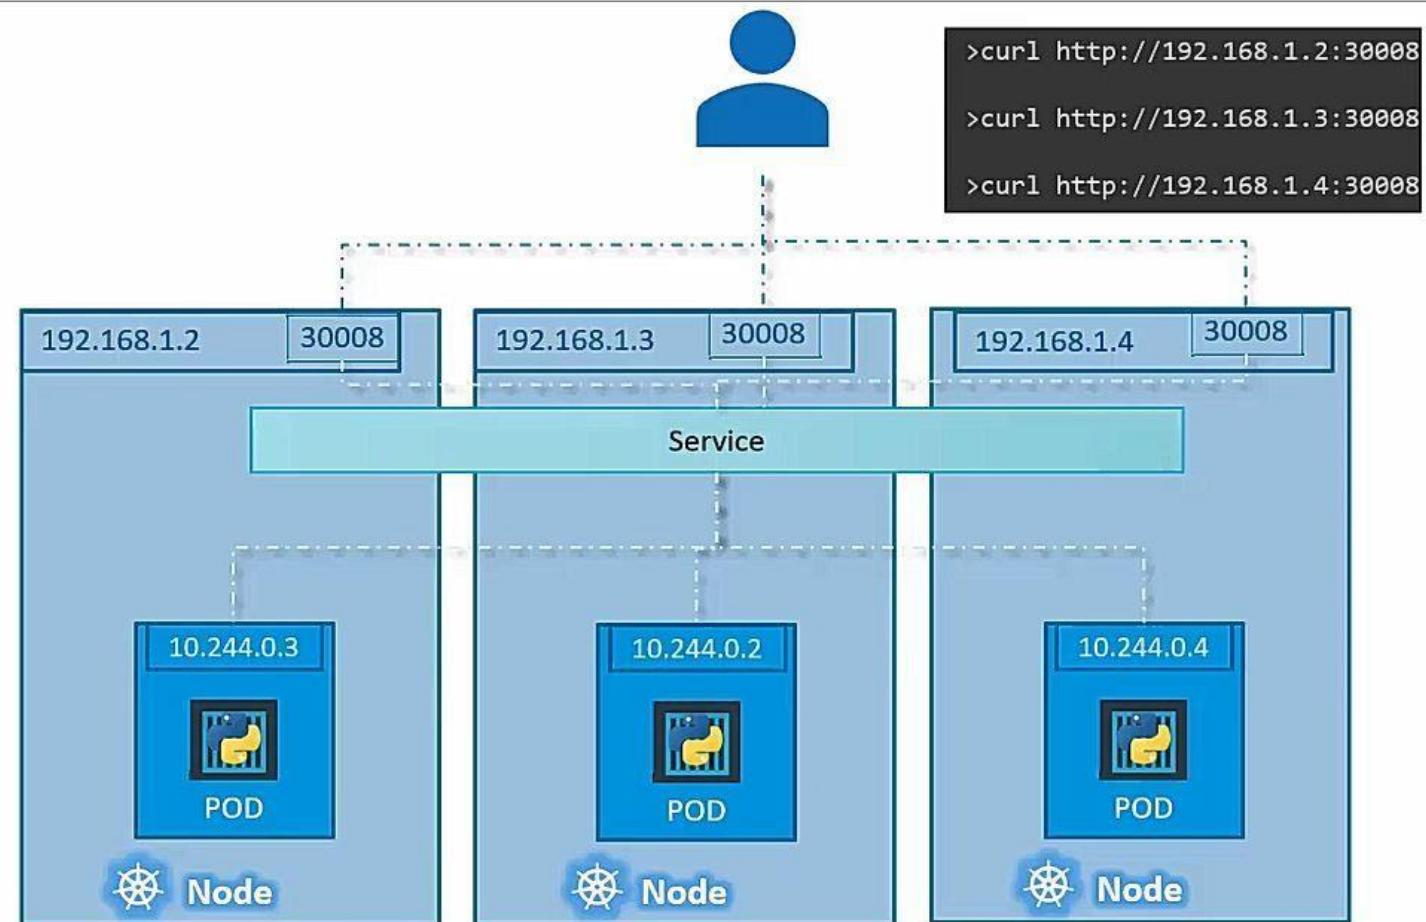
```
> curl http://192.168.1.2:30008
```

```
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
```

# Service - NodePort



# Service - NodePort



# LOAD BALANCER

- LoadBalancer: Exposes the Service externally using a cloud provider's load balancer. NodePort and ClusterIP Services, to which the external load balancer routes, are automatically created.

# MANUAL SCHEDULING

# SCHEDULER

A scheduler watches for newly created Pods that have no Node assigned. For every Pod that the scheduler discovers, the scheduler becomes responsible for finding the best Node for that Pod to run on. The scheduler reaches this placement decision taking into account the scheduling principles described below.

# SCHEDULING WITH KUBE SCHEDULER

kube-scheduler selects a node for the pod in a 2-step operation:

- Filtering
- Scoring

The filtering step finds the set of Nodes where it's feasible to schedule the Pod. For example, the `PodFitsResources` filter checks whether a candidate Node has enough available resource to meet a Pod's specific resource requests. After this step, the node list contains any suitable Nodes; often, there will be more than one. If the list is empty, that Pod isn't (yet) schedulable.

In the scoring step, the scheduler ranks the remaining nodes to choose the most suitable Pod placement. The scheduler assigns a score to each Node that survived filtering, basing this score on the active scoring rules.

Finally, kube-scheduler assigns the Pod to the Node with the highest ranking. If there is more than one node with equal scores, kube-scheduler selects one of these at random.

# MANUAL SCHEDULING

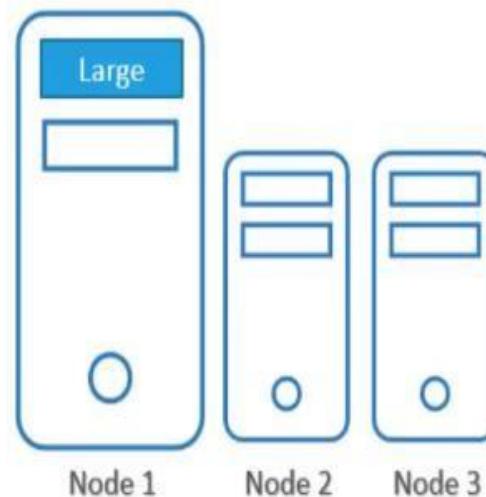
```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4      name: nginx
5      labels:
6          app: nginx
7  spec:
8      containers:
9          - name: nginx
10         image: nginx:1.7.9
11         ports:
12             - containerPort: 80
13         nodeName: node02
```

# NODE SELECTOR (LABEL ON NODES)

## Label Nodes

```
▶ kubectl label nodes <node-name> <label-key>=<label-value>
```

```
▶ kubectl label nodes node-1 size=Large
```



# NODE SELECTOR(LABEL ON PODS)

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: mypod
5 spec:
6   containers:
7     - name: data-processor
8       image: data-processor
9     nodeSelector:
10       size: Large
11
```

# RESOURCE ALLOCATION ON CONTAINER

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: simple-webapp
5   labels:
6     name: simple-webapp
7 spec:
8   containers:
9     - name: simple-webapp
10    image: simple-webapp
11    ports:
12      - containerPort: 8080
13    resources:
14      requests:
15        memory: "1Gi"
16        cpu: 1
17      limits:
18        memory: "2Gi"
19        cpu: 2
20
21
```

**NOTE:-** A container can not exceed 2 cpu limit if it tries to increase so it is known as throttle(limit) but it will continue working.

If container memory goes beyond 2Gi then it will terminated and create new if replica set is defined.

# LIMIT CPU AND MEMORY FOR CONTAINER

```
1 apiVersion: v1
2 kind: LimitRange
3 metadata:
4     name: limit-mem-cpu-per-container
5 spec:
6     limits:
7         - max:
8             cpu: "800m"
9             memory: "1Gi"
10        min:
11            cpu: "100m"
12            memory: "99Mi"
13        default:
14            cpu: "700m"
15            memory: "900Mi"
16        defaultRequest:
17            cpu: "110m"
18            memory: "111Mi"
19        type: Container
```

```
#kubectl create -f <filename>
```

```
#kubectl describe limitrange/limit-mem-cpu-per-container
```

# LIMITS MEMORY AND CPU FOR POD

```
1 apiVersion: v1
2 kind: LimitRange
3 metadata:
4     name: limit-mem-cpu-per-pod
5 spec:
6     limits:
7         - max:
8             cpu: "2"
9             memory: "2Gi"
10            type: Pod
11
12
13
```

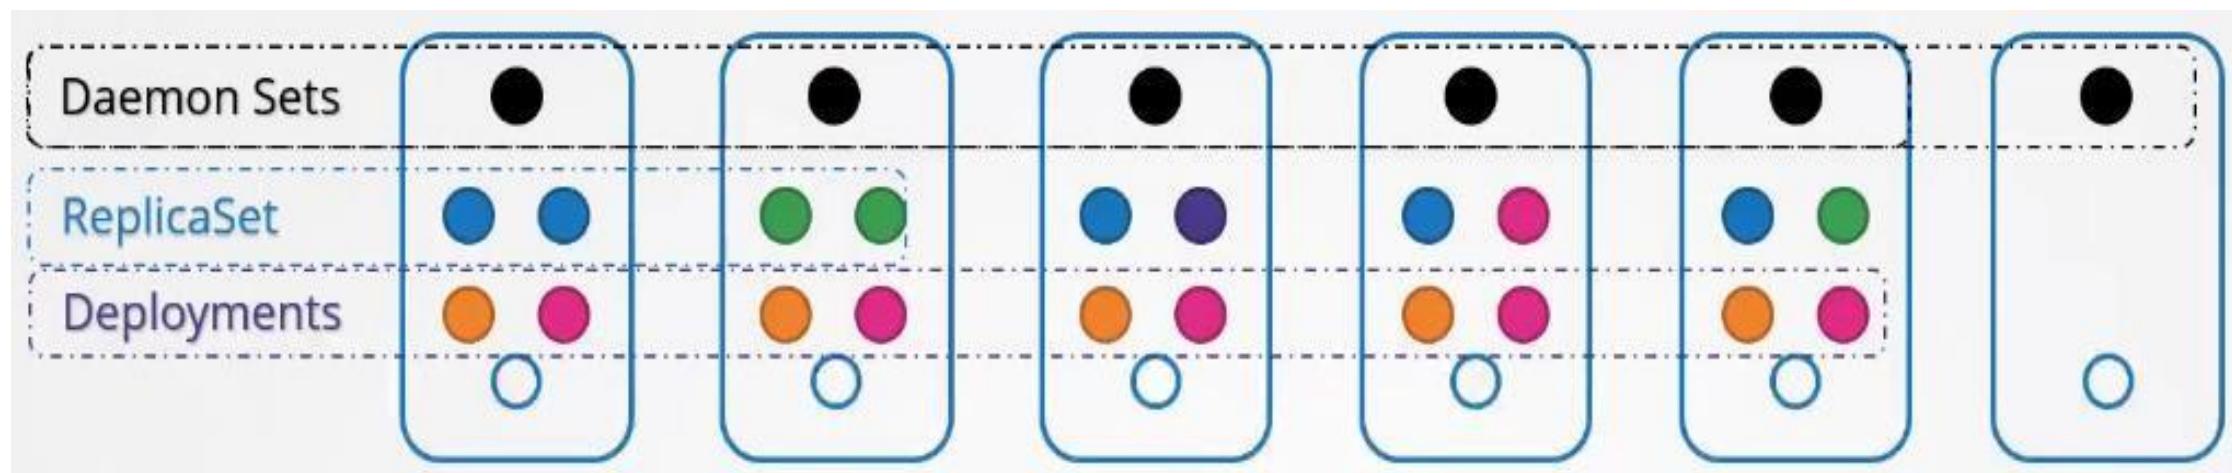
```
#kubectl create -f <filename>
#kubectl describe limitrange/limit-mem-cpu-per-pod
```

# Daemon Set

- A DaemonSet ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected. Deleting a DaemonSet will clean up the Pods it created.

Some typical uses of a DaemonSet are:

- running a cluster storage daemon, such as glusterd, ceph, on each node.
- running a logs collection daemon on every node, such as fluentd or filebeat.



# DAEMON SET

```
1 apiVersion: app/v1
2 kind: DaemonSet
3 metadata:
4   name: monitoring-daemon
5 spec:
6   selector:
7     matchLabels:
8       app: monitoring-agent
9   template:
10    metadata:
11      labels:
12        app: monitoring-agent
13    spec:
14      containers:
15        - name: monitoring
16          image: monitoring
17
```

# Deployment

- A Deployment provides declarative updates for pods and replicaset.
- You describe a desired state in a Deployment, and the Deployment Controller changes the actual state to the desired state at a controlled rate. You can define Deployments to create new ReplicaSets, or to remove existing Deployments and adopt all their resources with new Deployments.

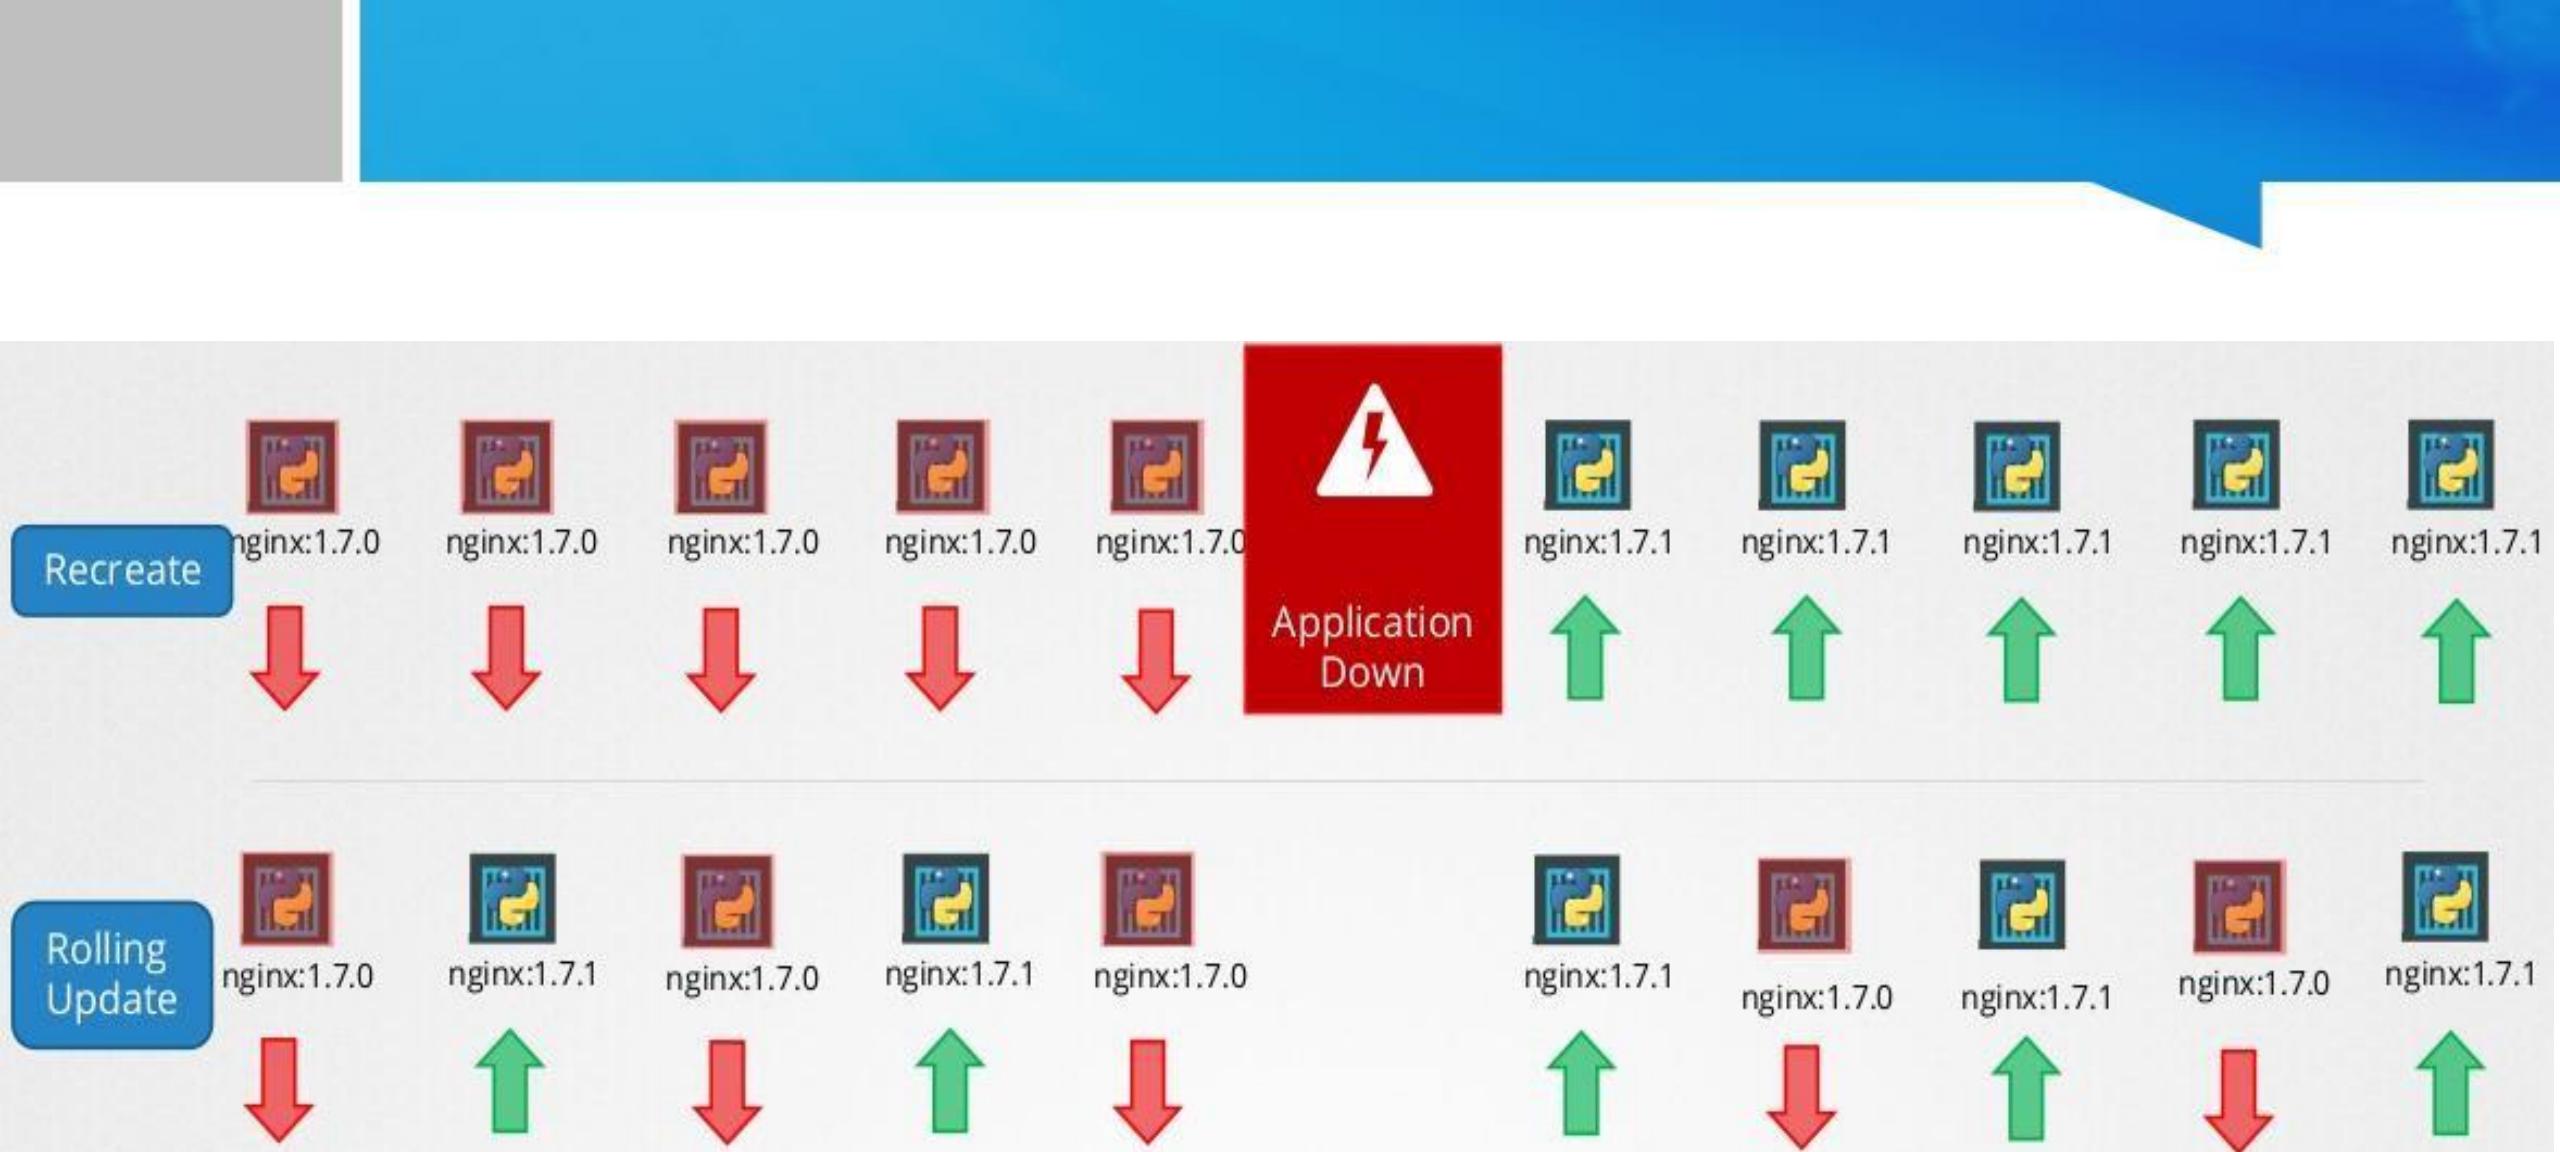
# Uses of Deployment

- Create a Deployment to rollout a ReplicaSet
- Declare the new state of the Pods
- Rollback to an earlier Deployment revision
- Scale up the Deployment to facilitate more load
- Pause the Deployment
- Use the status of the Deployment
- Clean up older ReplicaSets

# DEPLOYMENT STRATEGY

Two types of strategy->

- 1)Recreate -> in case of recreate lets say we have a pod and in pod there are 5 containers so it will first down all the containers then it will create new containers.  
so in a production environment it's not a good strategy because of there is downtime.
- 2)Rolling Update -> in case of rolling update it will down one container then update one then down the second container then create second.  
by default Kubernetes is using rolling update.



# DEPLOYMENT

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# ROLLOUT AND VERSIONING

- When we update any application so automatically it is called as revision so when we deploy any container so it takes as a revision 1 then when we update so revision2 and so on.

```
#kubectl rollout status deployment/myapp-deployment -> to check rollout status  
#kubectl rollout history deployment/myapp-deployment -> to check history of rollout
```

# DEPLOYMENT

To update manually->

```
#kubectl set image deployment/myapp-deployment nginx=nginx:1.7.1
```

To check deployment status and we can see strategy type->

```
#kubectl describe deployment myapp-deployment
```

- When you are doing upgrade so automatically it create a new replicaset.

```
#kubectl get rs
```

Now let's say we have some issue with update so we can rollout to previous version.->

```
#kubectl rollout undo deployment/myapp-dep
```

```
#kubectl get rs
```

# STORAGE

# Storage

## Volumes

- Exposed at Pod level and backend different ways
- emptyDir: ephemeral scratch directory that lives for the life of pod

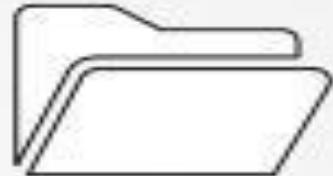
## Persistent Volumes

- Abstraction away from the storage provider(AWS, GCE)
- Have a lifecycle independent of any individual pod that uses it

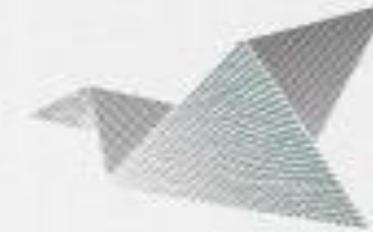
## Persistent Volume Claims

- Request for storage with specific details(size,access modes, etc.)

# TYPES OF KUBERNETES VOLUME



NFS



ceph

SCALEIO



# Persistent Volumes

- By default, containers write to ephemeral storage
- As a result, when a pod is terminated, all data written by its container is lost
- We can attach Persistent Volumes to pods which persist any data written to them

# emptyDir

An emptyDir volume is first created when a Pod is assigned to a Node, and exists as long as that Pod is running on that node

# Persistent Volume (emptyDir)

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pd
spec:
  containers:
    - image: k8s.gcr.io/test-webserver
      name: test-container
    volumeMounts:
      - mountPath: /cache
        name: cache-volume
  volumes:
    - name: cache-volume
      emptyDir: {}
```

# Persistent Volumes

- Persistent Volumes can be provisioned either statically or dynamically.
  - Static: A pre-provisioned pool of volumes such as iSCSI or Fiber Channel
  - Dynamic: Volumes are created on demand by calling a storage provider's API such as Amazon EBS
- Persistent Volumes have a lifecycle independent of the pods that use them

# Persistent Volume Claims

- Created by users to request a persistent volume
- Users can request various properties such as capacity and access modes (e.g. can be mounted once read/write or many times read-only)

# CREATING PERSISTENT VOLUME

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pv-voll
5 spec:
6   accessModes:
7     - ReadWriteOnce
8 capacity:
9   storage: 1Gi
10 hostPath:
11   path: "/tmp/data01"
12
13
14 ##three types of permissions-> 1) ReadOnlyMany 2) ReadWriteOnce 3) ReadWriteMany
15
16 #kubectl create -f <filename>
17 #kubectl get pv
18 #kubectl describe pv <pvname>
```

# CREATING PERSISTENT VOLUME CLAIM

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: myclaim
5 spec:
6   accessModes:
7     - ReadWriteOnce
8   resources:
9     requests:
10    storage: 500Mi
11
```

```
#kubectl create -f <filename>
#kubectl describe pvc
#kubectl describe pv <pvname>
```

# USING PV AND PVC IN POD

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: mypod
5   labels:
6     name: frontendhttp
7 spec:
8   containers:
9     - name: myweb
10    image: nginx
11    ports:
12      - containerPort: 80
13        name: "httpserver"
14    volumeMounts:
15      - mountPath: "/var/www/html"
16        name: myvol
17    volumes:
18      - name: myvol
19      persistentVolumeClaim:
20        claimName: myclaim ##give in pvc yml
```

# SECURITY

# Kubernetes Authentication

- Authentication mechanism handled by kubeapi server when we use kubectl command.

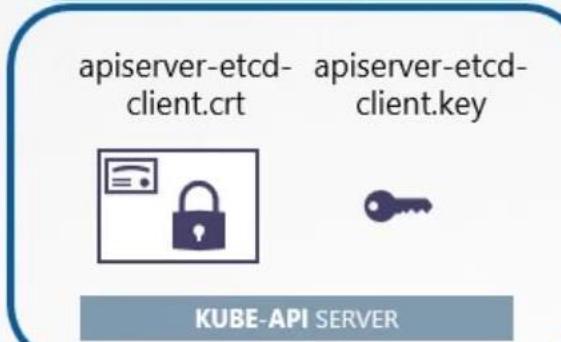
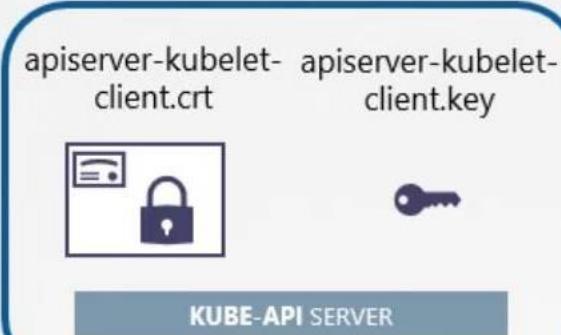
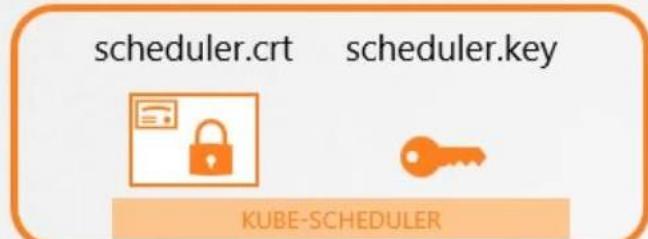
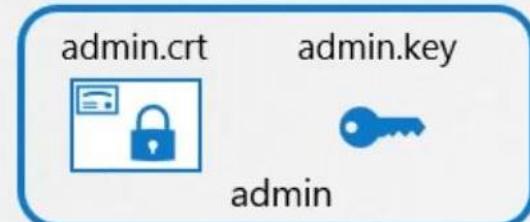
# Transport Layer Security (TLS)

# TLS

- Kubernetes provides a certificates.k8s.io API, which lets you provision TLS certificates signed by a Certificate Authority (CA) that you control. These CA and certificates can be used by your workloads to establish trust.



## Client Certificates for Clients



## Server Certificates for Servers



# Authorization

# Authorization

In Kubernetes, you must be authenticated (logged in) before your request can be authorized (granted permission to access).

Kubernetes authorizes API requests using the API server. It evaluates all of the request attributes against all policies and allows or denies the request. All parts of an API request must be allowed by some policy in order to proceed. This means that permissions are denied by default.

# Authorization Modes

- Node Based
  - grants permissions to kubelets based on the pods they are scheduled to run
- Attribute Based Access Control (ABAC)
  - Manual Manage Permissions in API
- Role Based Access Control (RBAC)
  - Roles and Role Bindings
- Webhook
  - manage third party authentication

# Kube API Server Yaml

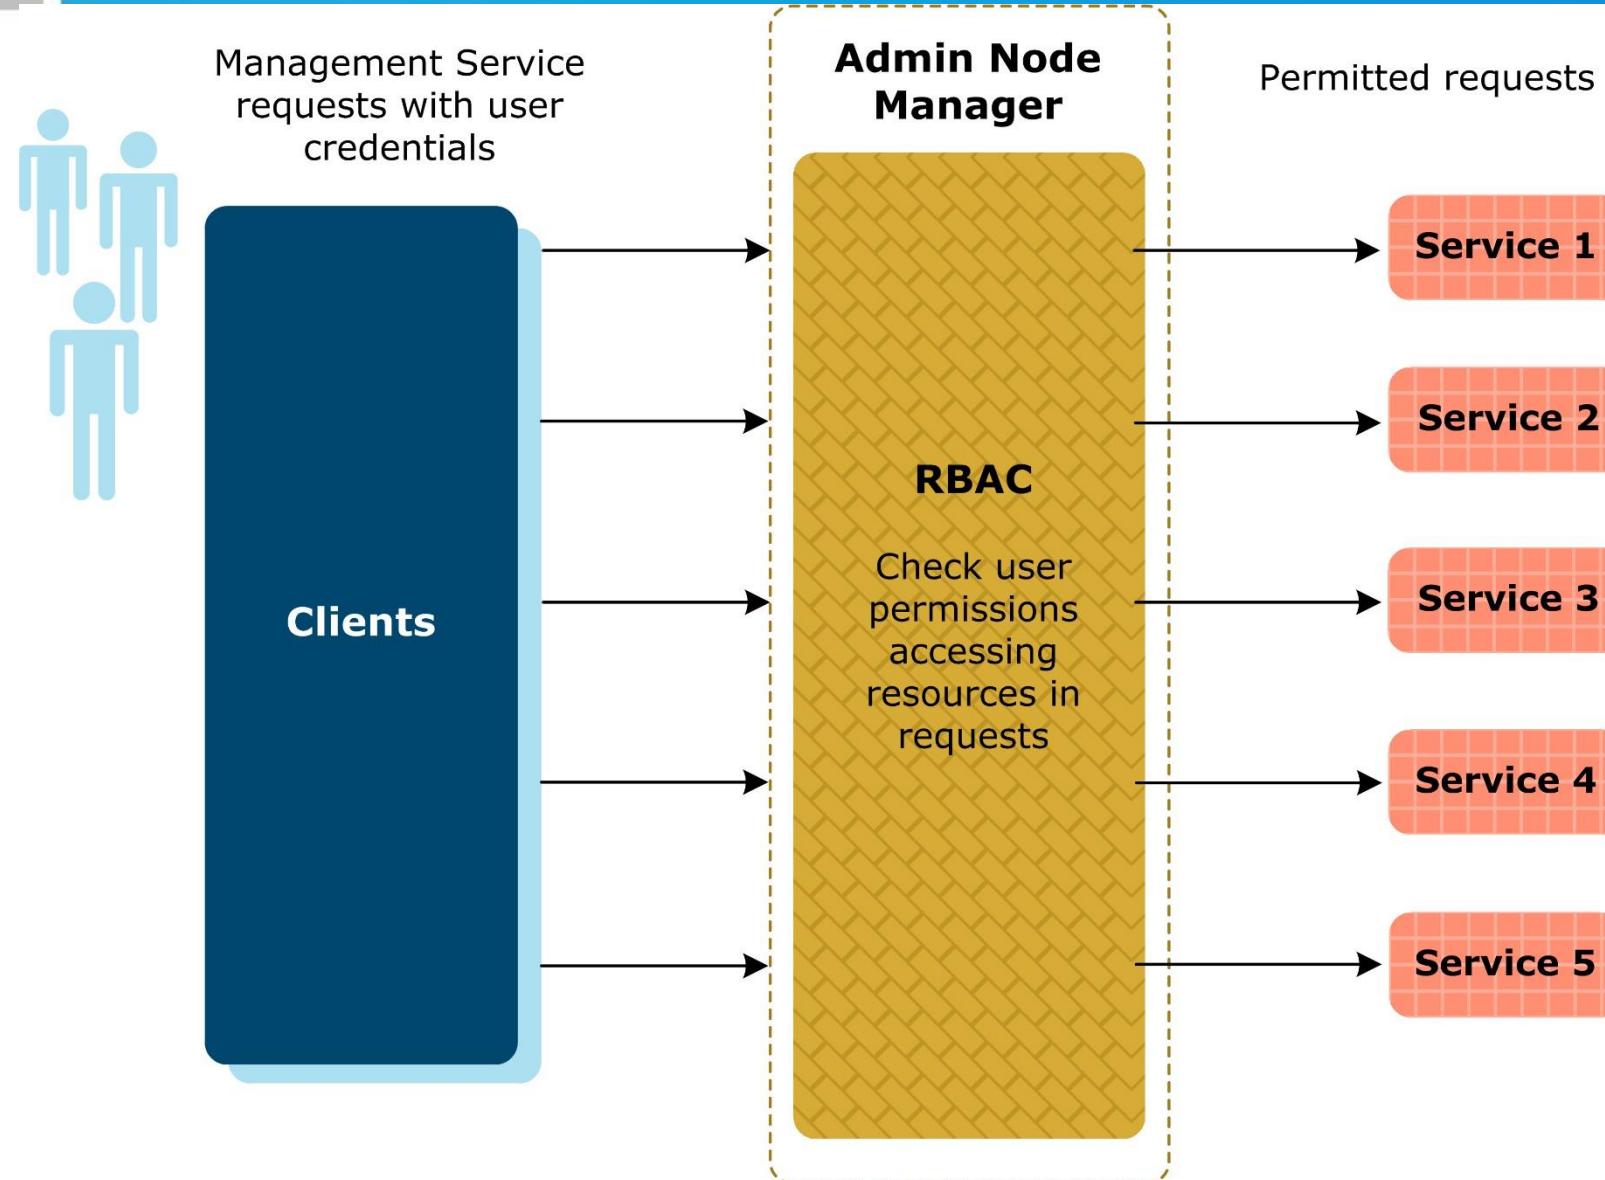
```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 172.25.230.143:6443
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-apiserver
    - --advertise-address=172.25.230.143
    - --allow-privileged=true
    - --authorization-mode=Node,RBAC
```

# Role-Based Access Control (RBAC) Authorization

# What is RBAC?

Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within your organization.

# RBAC



# Managing Roles with Imperative Commands

```
#kubectl create role pod-reader --verb=get --verb=list --verb=watch --  
resource=pods
```

```
#kubectl create role pod-reader --verb=get --resource=pods --resource-  
name=readablepod --resource-name=anotherpod
```

```
#kubectl create role foo --verb=get,list,watch --resource=replicasets.apps
```

# Verify

```
#kubectl get role  
#kubectl describe role <rolenam>
```

# Managing RoleBindings with Imperative

```
#kubectl create rolebinding bob-admin-binding --  
clusterrole=admin --user=bob --namespace=acme
```

```
#kubectl create rolebinding myapp-view-binding --  
clusterrole=view --serviceaccount=acme:myapp --  
namespace=acme
```

# Verify

```
#kubectl get rolebinding  
#kubectl describe rolebinding <name>
```

# Check Access

```
#kubectl auth can-i <verb> <resource>
```

For example:

```
#kubectl auth can-i create pods
```

```
#kubectl auth can-i create pods --as gaurav
```

# Note

**Roles and Rolebindings are restricted to namespace.**

**If we want to give access cluster wide then we will use clusterrole and clusterrolebinding.**

**To check role access->**

```
#kubectl api-resources --namespaced=true
```

**To check clusterrole access->**

```
#kubectl api-resources --namespaced=false
```

# Managing Cluster Role with Imperative

```
#kubectl create clusterrole pod-reader --verb=get,list,watch  
--resource=pods
```

```
#kubectl create clusterrole pod-reader --verb=get --  
resource=pods --resource-name=readablepod --resource-  
name=anotherpod
```

# Verify

```
#kubectl get clusterrole
```

```
#kubectl describe clusterrolebinding <name>
```

# Managing ClusterRoleBinding with Imperative

```
#kubectl create clusterrolebinding root-cluster-admin-binding --clusterrole=cluster-admin --user=root
```

```
#kubectl create clusterrolebinding myapp-view-binding --clusterrole=view --serviceaccount=acme:myapp
```

# SECURITY CONTEXT

Security context defines privilege and access control settings for a Pod or Container.

# SECURITY CONTEXT

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4     name: web-pod
5 spec:
6     securityContext:
7         runAsUser: 1000
8     containers:
9     - name: ubuntu
10        image: ubuntu
11        command: ["sleep", "3600"]
12
```

# SECURITY CONTEXT

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo-2
spec:
  securityContext:
    runAsUser: 1000
  containers:
  - name: sec-ctx-demo-2
    image: gcr.io/google-samples/node-hello:1.0
    securityContext:
      runAsUser: 2000
      allowPrivilegeEscalation: false
```

# SECURITY CONTEXT

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo-4
spec:
  containers:
  - name: sec-ctx-4
    image: gcr.io/google-samples/node-hello:1.0
    securityContext:
      capabilities:
        add: ["NET_ADMIN", "SYS_TIME"]
```

# Environment Variables

# TYPES OF ENVIRONMENT VARIABLES

- Plain key value
- ConfigMap
- Secrets

# PLAIN KEY VALUE

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: simple-webapp-color
5 spec:
6   containers:
7     - name: simple-webapp-color
8       image: simple-webapp-color
9       ports:
10      - containerPort: 8080
11   env:
12     - name: APP_COLOR
13       value: Pink
14 |
```

# CONFIG MAP

A ConfigMap is a dictionary of configuration settings. This dictionary consists of key-value pairs of strings. Kubernetes provides these values to your containers. Like with other dictionaries (maps, hashes, ...) the key lets you get and set the configuration value.

# CONFIG MAP

## STEP 1: CREATE CONFIG-MAP

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: app-config
5  data:
6    APP_COLOR: blue
7    APP_MODE: prod
8
9
10
```

# CONFIG MAP

## STEP 2: inject configmap in pod file

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4      name: myapp-pod
5      labels:
6          app: myapp
7  spec:
8      containers:
9          - name: nginx-container
10             image: nginx:1.7.1
11             ports:
12                 - containerPort: 8080
13             envFrom:
14                 - configMapRef:
15                     name: app-config
16
```

# SECRETS

Secrets can be defined as Kubernetes objects used to store sensitive data such as user name and passwords with encryption.

There are multiple ways of creating secrets in Kubernetes.

Creating from txt files.

Creating from yaml file.

# SECRETS

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: app-secret
5 data:
6   DB_Host: mysql
7   DB_User: root
8   DB_Password: password
9
10
11
12
```

in the above data we are specifying in a plain text so it's not a good idea so we can type it in a hash format. So we can convert in base64 format and type output in file.

```
#echo -n "mysql" | base64
bxIzcWw= --> output
```

# SECRETS

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: app-secret
5 data:
6   DB_Host: bxiZcWw=
7   DB_User: cb1Wx24x
8   DB_Password: gHTs7xD2
9
10
11
```

```
#kubectl get secrets
#kubectl describe secrets
#kubectl get secret app-secret -o yaml -> to show all information and encrypted password
#echo -n 'bxiZcWw=' | base64 --decode -> to decode
```

# SECRETS

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4      name: myapp-pod
5      labels:
6          app: myapp
7  spec:
8      containers:
9          - name: nginx-container
10         image: nginx:1.7.1
11         ports:
12             - containerPort: 8080
13         envFrom:
14             - secretRef:
15                 name: app-secret
16
```

# Logging and Monitoring

# What are the kinds of logs we want to capture?

- Container Logs
- Host OS Logs
- Control Plane Logs
- Event Messages

# Logging in Kubernetes and Containers

Helpful Command:

```
#kubectl logs <podname> -c <containername>
```

# Monitoring

- Free
  - Kube-state-metrics- collects metrics from Kubernetes API
  - Prometheus
  - metricbeat – standard metric monitoring that is Kubernetes aware
  - Ganglia
  - And more.....
- Third Party Services – Paid
  - Datadog
  - Honeycomb
  - Stackdriver with google monitoring
  - And more...

# Setup Prometheus Monitoring Tool



## Prometheus

## Monitoring

# Configuration

```
#git clone https://github.com/gauravkumar9130/kube-yaml.git
```

```
#cd kube-yaml
```

```
#kubectl create -f prometheus.yaml
```

```
#kubectl create -f kube-state-metrics-configs/.
```

```
#kubectl get pods -n monitoring
```

```
#kubectl get svc –n monitoring
```

Enable query history

Expression (press Shift+Enter for newlines)

Execute - insert metric at cursor - ↴

Graph Console

◀ Moment ▶

Element

no data

Value

Remove Graph

# NETWORKING

# DNS

- Kubernetes DNS schedules a DNS Pod and Service on the cluster, and configures the kubelets to tell individual containers to use the DNS Service's IP to resolve DNS names.
- Every Service defined in the cluster (including the DNS server itself) is assigned a DNS name. By default, a client Pod's DNS search list will include the Pod's own namespace and the cluster's default domain.

# CUSTOM DNS FOR CONTAINER

```
apiVersion: v1
kind: Pod
metadata:
  name: hostaliases-pod
spec:
  restartPolicy: Never
  hostAliases:
  - ip: "127.0.0.1"
    hostnames:
    - "foo.local"
    - "bar.local"
  - ip: "10.1.2.3"
    hostnames:
    - "foo.remote"
    - "bar.remote"
  containers:
  - name: cat-hosts
    image: busybox
    command:
    - cat
    args:
    - "/etc/hosts"
```

# NAMESPACE

- Namespace provides an additional qualification to a resource name. This is helpful when multiple teams are using the same cluster and there is a potential of name collision. It can be as a virtual wall between multiple clusters.
- It offers an isolation for process interaction within OS.
- It limits the visibility that a process has on other process, network filesystem and userID component.
- Process from the containers or the host processes are not directly accessible from within this container process.

# FUNCTIONALITY OF NAMESPACE

- Namespaces help pod-to-pod communication using the same namespace.
- Namespaces are virtual clusters that can sit on top of the same physical cluster.
- They provide logical separation between the teams and their environments.

## Namespace

### Namespace 1

**Service**

IP address: 80

Pod 1

Pod 2

**ReplicationSet/  
Replication Controller**

### Namespace 2

**Service**

IP address: 80

Pod 1

Pod 2

**ReplicationSet/  
Replication Controller**

# CREATE NAMESPACE

```
#vim name.yml
```

```
1 apiVersion: v1
2 kind: Namespace
3 metadata:
4   name: dev
5 
```

```
[root@kmaster ~]# kubectl create -f name.yml
namespace/myname created
```

```
#kubectl create namespace dev
```

# TO ADD PODS IN NAMESPACE

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4     name: myapp-pod
5     namespace: dev
6     labels:
7         app: myapp
8         type: front-end
9 spec:
10    containers:
11        - name: nginx-container
12          image: nginx
13
14
```

# CONTROL THE NAMESPACES

```
# kubectl create -f namespace.yml  
# kubectl get namespace  
# kubectl get namespace <Namespace name>  
# kubectl describe namespace <Namespace name>  
# kubectl delete namespace <Namespace name>
```

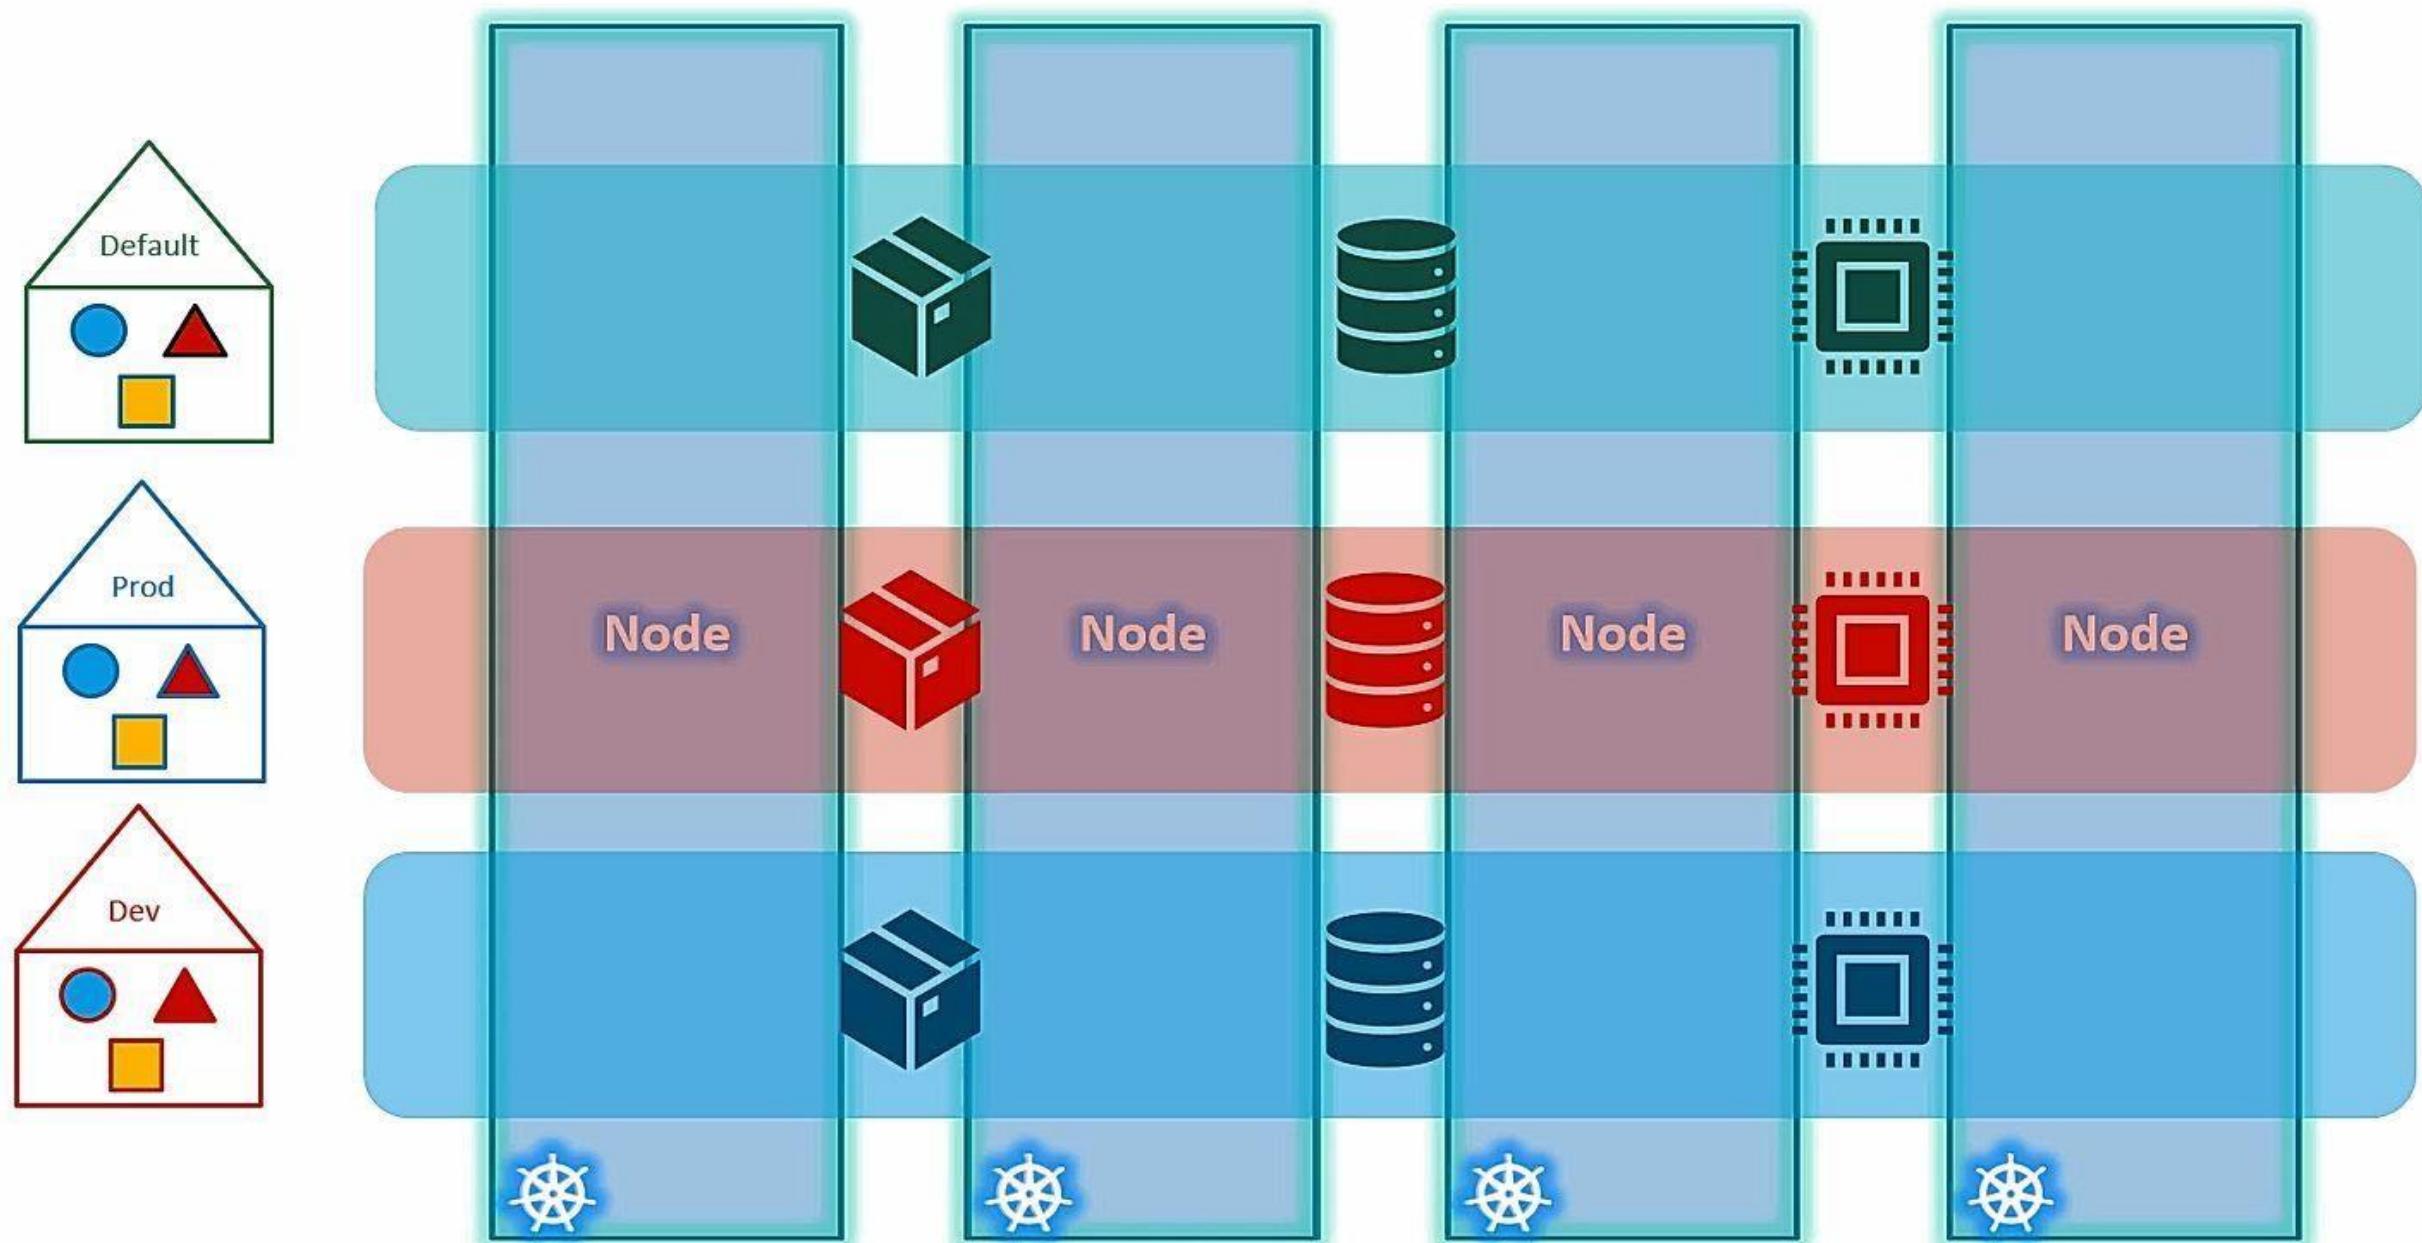
To switch from default namespace to another namespace:

```
# kubectl config set-context $(kubectl config current-context) --namespace=dev
```

To view pods from all namespaces

```
# kubectl get pods --all-namespaces
```

# Namespace – Resource Limits



# RESOURCE QUOTA IN NAMESPACE

```
1  apiVersion: v1
2  kind: ResourceQuota
3  metadata:
4      name: compute-quota
5      namespace: dev
6  spec:
7      hard:
8          pods: "10"
9          requests.cpu: "4"
10         requests.memory: 5Gi
11         limits.cpu: "10"
12         limits.memory: 10Gi
13
14
15 #kubectl create -f compute-quota.yml
16
```

# Create Pod In given Namespace

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: simple-webapp
5   labels:
6     name: simple-webapp
7 spec:
8   containers:
9     - name: simple-webapp
10    image: simple-webapp
11    ports:
12      - containerPort: 8080
13    resources:
14      requests:
15        memory: "1Gi"
16        cpu: 1
17      limits:
18        memory: "2Gi"
19        cpu: 2
20
21
```

# CONTAINER NETWORK INTERFACE

- Kubernetes does not provide any default network implementation, rather it only defines the model and leaves to other tools to implement it. We can use **Calico** Solution.

# NETWORK SOLUTION

We have popular Network Solution->

- Calico (<https://www.tigera.io/blog/kubernetes-networking-with-calico/>)

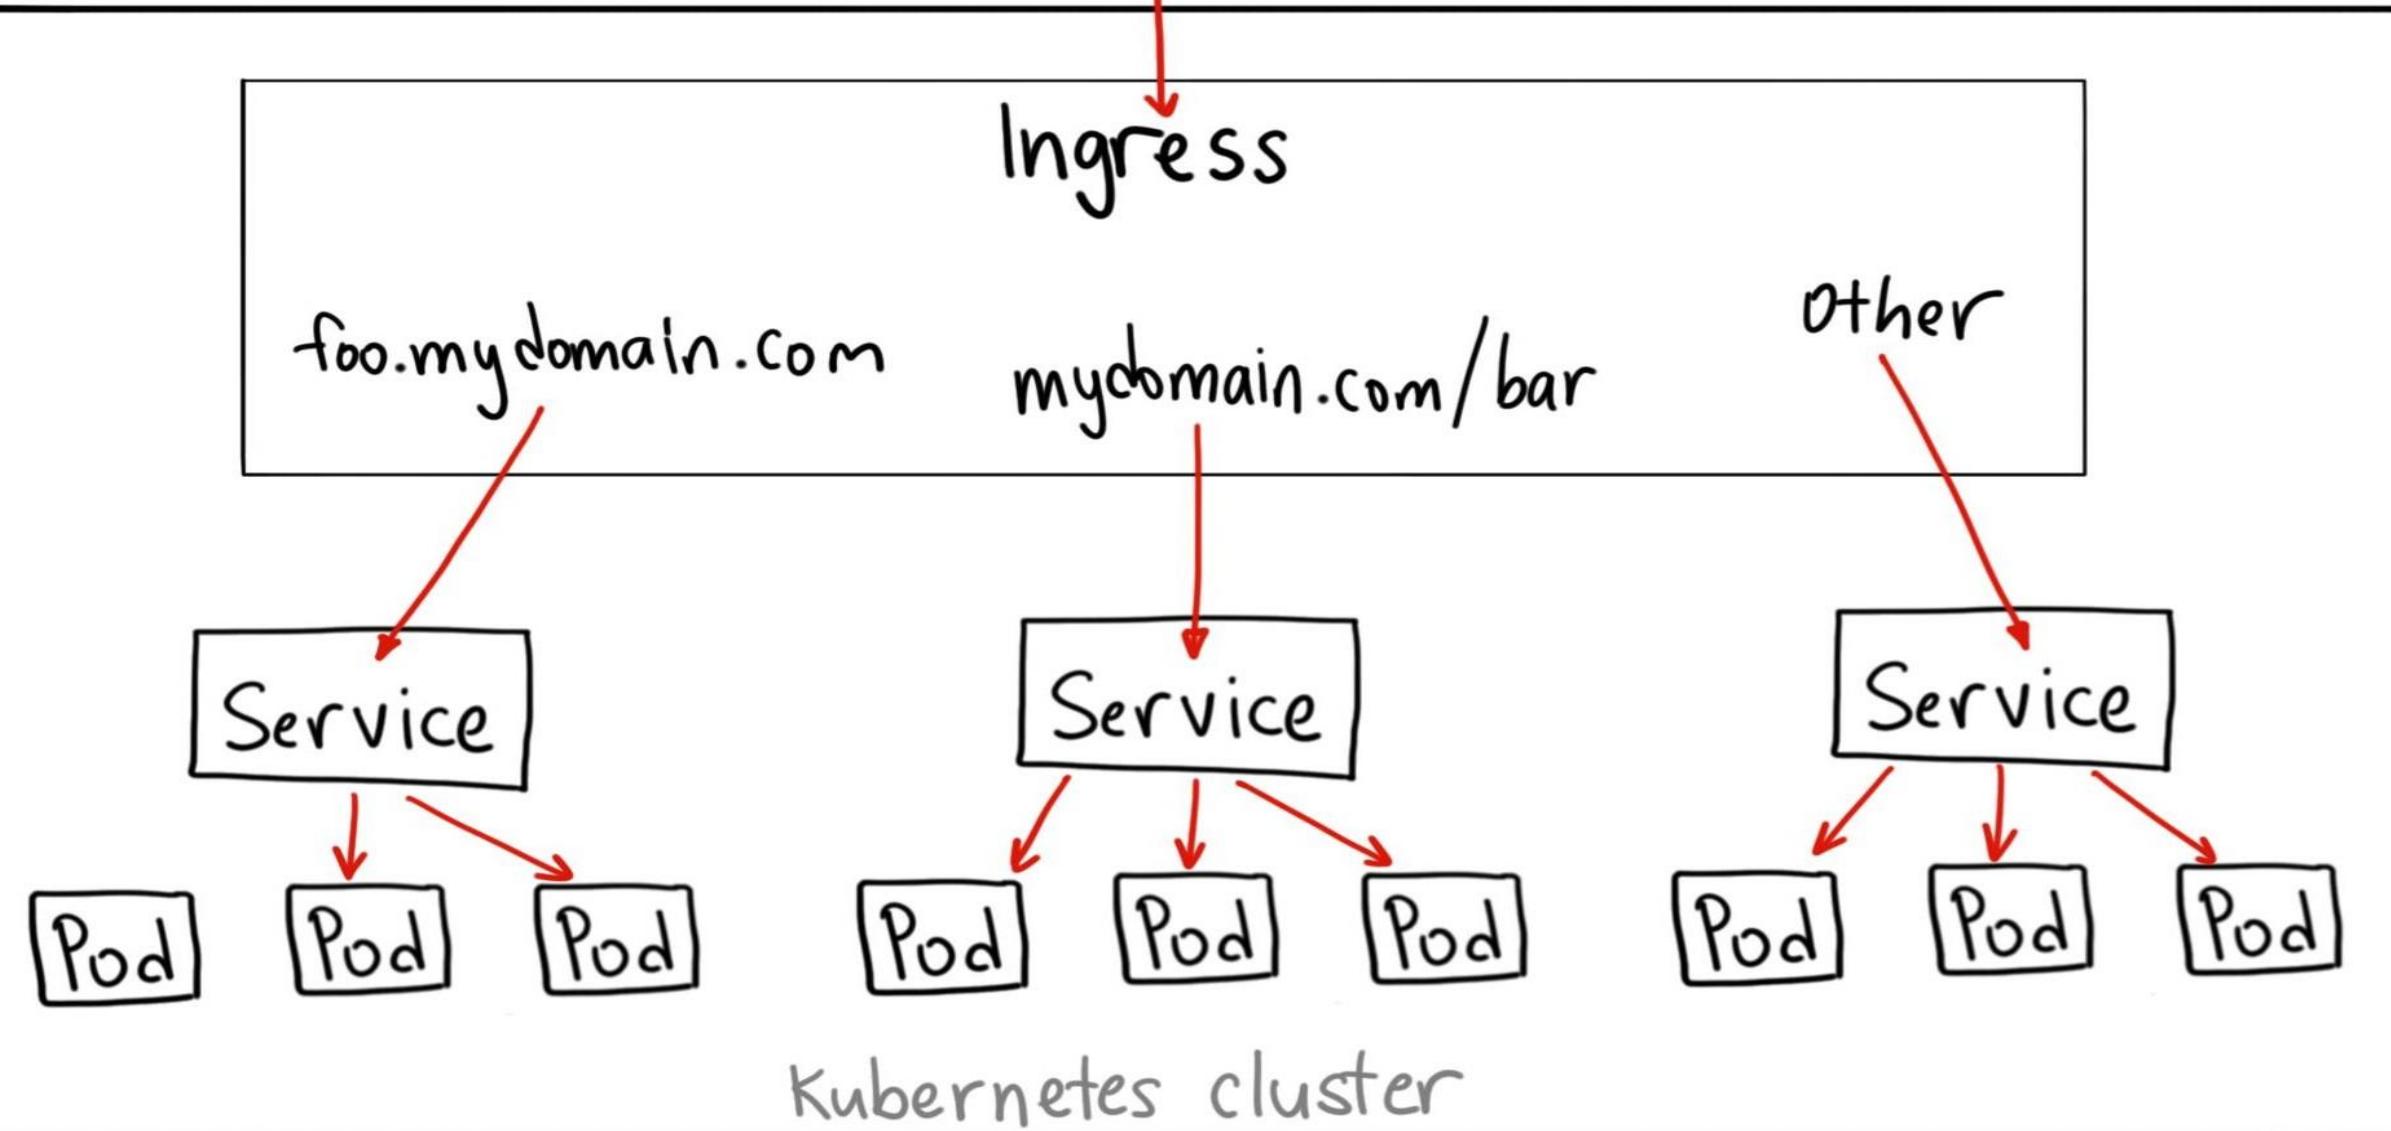
<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

# CALICO

- Calico is a networking and network policy provider. Calico supports a flexible set of networking options so you can choose the most efficient option for your situation.
- By default, Calico uses any ipaddress other than your cluster ip, as the Pod network CIDR, though this can be configured in the calico.yaml file. For Calico to work correctly, you need to pass this same CIDR to the kubeadm init command using the **--pod-network-cidr** flag or via kubeadm's configuration.
- `kubectl apply -f https://docs.projectcalico.org/v3.14/manifests/calico.yaml`

# INGRESS

- In Kubernetes, an Ingress is an object that allows access to your Kubernetes services from outside the Kubernetes cluster. You configure access by creating a collection of rules that define which inbound connections reach which services.



# HOW TO DEPLOY INGRESS

- `#git clone https://github.com/nginxinc/kubernetes-ingress.git`
- `#cd kubernetes-ingress`
- `#cd deployments`

Create a namespace and service account->

- `#kubectl apply -f common/ns-and-sa.yaml`

Create a secret with a tls certificate and a key for the default server in nginx:-

- `#kubectl apply -f common/default-server-secret.yaml`

Set the behaviour of nginx-config (we can change)->

- `#kubectl apply -f common/nginx-config.yaml`

Configure RBAC->

- `#kubectl apply -f rbac/rbac.yaml`

# HOW TO DEPLOY INGRESS

Deploy the ingress controller->

We have two options ->

- 1) deployment
- 2) daemon set -> we will use because it will be available in every container

```
#kubectl apply -f daemon-set/nginx-ingress.yaml
```

# Verify

- `#kubectl get namespace`
- `#kubectl get all -n nginx-ingress`
- `#kubectl get all -n nginx-ingress`
- `#kubectl get pods -n nginx-ingress -o wide`

# **Host Based Ingress (Hotel Application)**

# Application

- we are creating a website for :

**hotel.example.com**

With using hotel application and service

# Application(HostBased)

```
#mkdir ingress-hostbased
```

Creating Namespace:

```
#vim create_namespace.yml
apiVersion: v1
kind: Namespace
metadata:
  name: hotel
```

# Deploying Hotel Application

```
#vim app_deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hotel
  namespace: hotel
spec:
  replicas: 2
  selector:
    matchLabels:
      app: hotel
  template:
    metadata:
      labels:
        app: hotel
    spec:
      containers:
        - name: hotel
          image: nginxdemos/hello:plain-text
          ports:
            - containerPort: 80
```

# Creating Service for Application

```
#vim service_for_hotel.yml
apiVersion: v1
kind: Service
metadata:
  name: hotel-svc
  namespace: hotel
spec:
  ports:
  - port: 80
    targetPort: 80
    protocol: TCP
    name: http
  selector:
    app: hotel
```

# Create Ingress Rule(Host Based)

```
#vim ingress_rule.yml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: hotel-ingress
  namespace: hotel
spec:
  rules:
  - host: hotel.example.com
    http:
      paths:
      - path: /
        backend:
          serviceName: hotel-svc
          servicePort: 80
```

# Verify

- `#kubectl get all --namespace hotel`
- `#kubectl describe ingress hotel-ingress --namespace hotel`

**Add entry in host file because of we are not using any loadbalancer:**

`#vim /etc/hosts`

`<ipofworker> cafe.example.com`

# **Path Based Ingress (Cafe Application)**

# Application

- we are creating a website for : cafe.example.com
- which is having micro-services running
  - café.example.com/tea -> tea-app
  - café.example.com/coffee -> coffee-app

# Create Namespace

```
#vim create_namespace.yml
apiVersion: v1
kind: Namespace
metadata:
  name: cafe
```

# Create Coffee-App

```
#vim app_deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: coffee
  namespace: cafe
spec:
  replicas: 2
  selector:
    matchLabels:
      app: coffee
  template:
    metadata:
      labels:
        app: coffee
    spec:
      containers:
        - name: coffee
          image: nginxdemos/hello:plain-text
          ports:
            - containerPort: 80
```

# Create Tea-App

```
#vim app_deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: tea
  namespace: cafe
spec:
  replicas: 2
  selector:
    matchLabels:
      app: tea
  template:
    metadata:
      labels:
        app: tea
    spec:
      containers:
        - name: tea
          image: nginxdemos/hello:plain-text
          ports:
            - containerPort: 80
```

# Service for Coffee App

```
#vim service_for_coffee.yml
apiVersion: v1
kind: Service
metadata:
  name: coffee-svc
  namespace: cafe
spec:
  ports:
  - port: 80
    targetPort: 80
    protocol: TCP
    name: http
  selector:
    app: coffee
```

# Service for Tea App

```
#vim service_for_tea.yml
apiVersion: v1
kind: Service
metadata:
  name: tea-svc
  namespace: cafe
spec:
  ports:
  - port: 80
    targetPort: 80
    protocol: TCP
    name: http
  selector:
    app: tea
```

# Ingress Rule

```
#vim ingress_rule.yml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: cafe-ingress
  namespace: cafe
spec:
  rules:
    - host: cafe.example.com
      http:
        paths:
          - path: /tea
            backend:
              serviceName: tea-svc
              servicePort: 80
          - path: /coffee
            serviceName: coffee-svc
            servicePort: 80
```

# Verify Application

```
#kubectl get ingress --namespace cafe  
#kubectl describe ingress cafe-ingress --namespace cafe (show the Rules: )  
#kubectl get pods -o wide -n café
```

**Add entry in host file because of we are not using any loadbalancer**

**#vim /etc/hosts**

- ip of worker cafe.example.com**

# Load Balancer

# WHY?

- Kubernetes does not offer an implementation of network load-balancers (Services of type LoadBalancer) for bare metal clusters.
- LoadBalancers will remain in the “pending” state indefinitely when created.

# METAL LOAD BALANCER

- MetallB hooks into your Kubernetes cluster, and provides a network load-balancer implementation. In short, it allows you to create Kubernetes services of type “LoadBalancer” in clusters.
- It has two features that work together to provide this service: address allocation, and external announcement.

# Prerequisites

```
#kubectl edit configmap -n kube-system kube-proxy
```

```
apiVersion: kubeproxy.config.k8s.io/v1alpha1
kind: KubeProxyConfiguration
mode: "ipvs"
ipvs:
  strictARP: true
```

# Installation By Manifests



```
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.9.3/manifests/namespace.yaml
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.9.3/manifests/metallb.yaml
# On first install only
kubectl create secret generic -n metallb-system memberlist --from-literal=secretkey="$(openssl rand -base64
128)"
```

<https://metallb.universe.tf/installation/>

# Create ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.1.240-192.168.1.250
```

<https://metallb.universe.tf/configuration/>

# LET'S DEPLOY A APPLICATION

```
[root@master abc]# cat deploy.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 4
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      name: nginx-deployment
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx-deployment
          image: nginx
          ports:
            - containerPort: 80
```

# SERVICE FOR APPLICATION

```
[root@master abc]# cat service.yml
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: nginx
  type: LoadBalancer
[root@master abc]# ■
```

# Verify

```
[root@master abc]# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	10h
nginx	LoadBalancer	10.98.253.120	172.25.230.240	80:31739/TCP	13m
[root@master abc]# █					

# Deploying PHP Guestbook application with Redis

# Application

- A single-instance Redis master to store guestbook entries
- Multiple replicated Redis instances to serve reads
- Multiple web frontend instances

## Reference:

<https://kubernetes.io/docs/tutorials/stateless-application/guestbook/>



# THANK YOU